

Final Report

March 12, 2021

1 MGTA415 Group Project

1.1 Reviews Analysis And Recommendation System For E-commerce Retails

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import re
import os
import time

from unidecode import unidecode
import contractions
import itertools
from collections import Counter

from nltk import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
import nltk
from gensim.models import word2vec
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.feature_extraction.text import (
    CountVectorizer,
    TfidfTransformer,
    TfidfVectorizer,
)
from sklearn import metrics
from sklearn.metrics import f1_score
import xgboost
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
from keras.utils.np_utils import to_categorical
from keras.callbacks import EarlyStopping
from keras.layers import Dense
from keras.models import Sequential
from sklearn.preprocessing import LabelEncoder
```

```

from math import sqrt
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from math import sqrt
import warnings
from wordcloud import WordCloud
warnings.filterwarnings("ignore")

```

1.2 Load data

```

[2]: reviews=pd.read_csv("data/Womens Clothing E-Commerce Reviews.csv")
reviews["Rating"]=reviews["Rating"].astype("category")
reviews=reviews[reviews['Review Text'].notnull()]
reviews.head(2)

```

```

[2]:   Unnamed: 0  Clothing ID  Age Title  \
0           0           767   33   NaN
1           1          1080   34   NaN

                                Review Text Rating  Recommended IND  \
0  Absolutely wonderful - silky and sexy and comf...         4             1
1  Love this dress!  it's sooo pretty.  i happene...         5             1

Positive Feedback Count Division Name Department Name Class Name
0                      0      Initmates      Intimate  Intimates
1                      4       General      Dresses   Dresses

```

```

[3]: reviews.Rating.value_counts()

```

```

[3]: 5    12540
4     4908
3     2823
2     1549
1       821
Name: Rating, dtype: int64

```

1.3 Data preprocessing

1.3.1 2.1 initial method

```

[4]: def text_clean(df,colname):

    nltk.download('stopwords')
    stop_words = stopwords.words('english')
    # stop_words.extend(add_stop_words)
    stop_words = set(stop_words)

```

```

ps = PorterStemmer()

preprocessed_t_sentences = []

for i, row in df.iterrows():
    # Expand contractions
    sent= contractions.fix(str(row[colname]))
    # Remove HTML tags
    sent=re.sub("<.*?>", "", sent)
    # Remove numbers and puncs
    sent=re.sub("\\r\\n", " ", sent)
    sent=re.sub("[^a-zA-Z\\s]", " ", sent.lower())
    sent=re.sub("\\s+", " ", sent)
    sent=re.sub(r'([\\w])\\1\\1+', r'\\1', sent)
#     sent_t=re.sub("(\\W|\\d)", " ", sent_t)
    words_list = sent.strip().split()
    #lowercasing, standardized english, remove stop words
    filtered_words = [ps.stem(unidecode(word)) for word in words_list if
    ↪word not in stop_words and len(word) != 1 and ps.stem(word) not in
    ↪stop_words]
#     preprocessed_t_sentences.append(" ".join(filtered_t_words))
#     preprocessed_t_sentences.append(filtered_t_words)
    preprocessed_t_sentences.append(" ".join(filtered_words))
    df[colname+"_processed"]=preprocessed_t_sentences
    return df

```

```

[5]: df=text_clean(reviews,"Review Text")
df.head()

```

```

[nltk_data] Downloading package stopwords to
[nltk_data]   D:\WispZ\AppData\Roaming\nltk_data...
[nltk_data]   Package stopwords is already up-to-date!

```

```

[5]:
   Unnamed: 0  Clothing ID  Age  Title \
0           0           767   33      NaN
1           1          1080   34      NaN
2           2          1077   60  Some major design flaws
3           3          1049   50    My favorite buy!
4           4           847   47    Flattering shirt

```

```

                                Review Text  Rating  Recommended  IND  \
0  Absolutely wonderful - silky and sexy and comf...      4          1
1  Love this dress!  it's sooo pretty.  i happene...      5          1
2  I had such high hopes for this dress and reall...      3          0
3  I love, love, love this jumpsuit. it's fun, fl...      5          1
4  This shirt is very flattering to all due to th...      5          1

```

	Positive Feedback Count	Division Name	Department Name	Class Name	\
0	0	Initmates	Intimate	Intimates	
1	4	General	Dresses	Dresses	
2	0	General	Dresses	Dresses	
3	0	General Petite	Bottoms	Pants	
4	6	General	Tops	Blouses	

```

                                Review Text_processed
0                                absolut wonder silki sexi comfort
1  love dress pretti happen find store glad bc ne...
2  high hope dress realli want work initi order p...
3  love love love jumpsuit fun flirti fabul everi...
4  shirt flatter due adjust front tie perfect len...

```

1.3.2 2.2 filter low frequency words

```

[6]: tagged_data=[d for d in df["Review Text_processed"].apply(lambda x:x.split("
→"))]
word_counts=Counter(itertools.chain(*tagged_data))
freq_dict={pair[0]:pair[1] for pair in word_counts.most_common()}
df_freq=pd.DataFrame(freq_dict.items()).rename(columns={0:"word",1:"freq"})
low_frequent_words=df_freq[df_freq.freq<5].word.tolist()
len(low_frequent_words)

```

[6]: 5768

```

[7]: def text_clean_freq(df,colname,low_freq_list):

    nltk.download('stopwords')
    stop_words = stopwords.words('english')
    stop_words.extend(low_freq_list)
    stop_words = set(stop_words)
    ps = PorterStemmer()

    preprocessed_t_sentences = []

    for i, row in df.iterrows():
        # Expand contractions
        sent= contractions.fix(str(row[colname]))
        # Remove HTML tags
        sent=re.sub("<.*?>", "",sent)
        # Remove numbers and puncs
        sent=re.sub("\r\n", " ",sent)
        sent=re.sub("[^a-zA-Z\s]", " ",sent.lower())
        sent=re.sub("\s+", " ",sent)
        sent=re.sub(r'([\w])\1\1+', r'\1', sent)
    #     sent_t=re.sub("(\W|\d)", " ",sent_t)

```

```

words_list = sent.strip().split()
#lowercasing, standardized english, remove stop words
filtered_words = [ps.stem(unidecode(word)) for word in words_list if
↪word not in stop_words and len(word) != 1 and ps.stem(word) not in
↪stop_words]
#         preprocessed_t_sentences.append(" ".join(filtered_t_words))
#         preprocessed_t_sentences.append(filtered_t_words)
preprocessed_t_sentences.append(" ".join(filtered_words))
df[colname+"_processed"]=preprocessed_t_sentences
return df

```

```

[8]: df_lf=text_clean_freq(reviews,"Review Text",low_frequent_words)
df_lf.head()

```

```

[nltk_data] Downloading package stopwords to
[nltk_data]   D:\WispZ\AppData\Roaming\nltk_data...
[nltk_data]   Package stopwords is already up-to-date!

```

```

[8]: Unnamed: 0  Clothing ID  Age  Title \
0          0          767   33      NaN
1          1         1080   34      NaN
2          2         1077   60  Some major design flaws
3          3         1049   50    My favorite buy!
4          4          847   47    Flattering shirt

```

```

                                Review Text Rating  Recommended IND  \
0  Absolutely wonderful - silky and sexy and comf...      4          1
1  Love this dress! it's sooo pretty. i happene...      5          1
2  I had such high hopes for this dress and reall...      3          0
3  I love, love, love this jumpsuit. it's fun, fl...      5          1
4  This shirt is very flattering to all due to th...      5          1

```

```

Positive Feedback Count  Division Name Department Name Class Name  \
0          0          Initmates      Intimate  Intimates
1          4          General      Dresses  Dresses
2          0          General      Dresses  Dresses
3          0  General Petite      Bottoms  Pants
4          6          General      Tops  Blouses

```

```

                                Review Text_processed
0          absolut wonder silki comfort
1  love dress pretti happen find store glad bc ne...
2  high hope dress realli want work initi order p...
3  love love love jumpsuit fun flirti fabul everi...
4  shirt flatter due adjust front tie perfect len...


```

1.4 Exploratory Data Analysis


```
[9]: sns.set_style("whitegrid")
test = df_lf
test = test[~test['Review Text'].isnull()]
```

```
[10]: plt.figure(figsize = (20,6))
ax = sns.countplot(x = 'Age', data = reviews,palette = 'viridis')

plt.show()
figure = ax.get_figure()
figure.savefig('age.png', dpi=500)
```



```
[11]: plt.figure(figsize = (8,6))
y = test['Review Text_processed']
X = test.drop(columns = 'Review Text_processed')
X = X.drop(columns = 'Unnamed: 0')
sns.heatmap(X.corr(), annot = True )
plt.xticks(rotation = 45)
plt.yticks(rotation = 45)
plt.savefig('heatmap.png', dpi=500)
```



1.4.1 WordCloud

```
[13]: nltk.download('brown')
```

```
[nltk_data] Downloading package brown to /home/jovyan/nltk_data...
[nltk_data] Package brown is already up-to-date!
```

```
[13]: True
```

```
[22]: from textblob import *
nltk.download('averaged_perceptron_tagger')
comment_words = ''
stop_words = set(stopwords.words("english"))
for val in reviews['Review Text']:
    # typecaste each val to string
    val = str(val)

    # split the value
```

```

tokens = val.split()

# Converts each token into lowercase
for i in range(len(tokens)):
    tokens[i] = tokens[i].lower()

comment_words += " ".join(tokens)+" "

def extract_NN(sent):
    grammar = r"""
    NBAR:
        # Nouns and Adjectives, terminated with Nouns
        {<NN.*>*<NN.*>}

    NP:
        {<NBAR>}
        # Above, connected with in/of/etc...
        {<NBAR><IN><NBAR>}
    """
    chunker = nltk.RegexpParser(grammar)
    ne = ''
    chunk = chunker.parse(nltk.pos_tag(nltk.word_tokenize(sent)))
    for tree in chunk.subtrees(filter=lambda t: t.label() == 'NP'):
        ne += ' '.join([child[0] for child in tree.leaves()]) + " "
    return ne

words = extract_NN(comment_words)
wordcloud = WordCloud(width = 800, height = 800,
                       background_color = 'white',
                       stopwords = stop_words,
                       min_font_size = 16).generate(words)

# plot the WordCloud image
plt.figure(figsize = (8, 8), facecolor = None)
ax = plt.imshow(wordcloud)
plt.axis("off")
plt.tight_layout(pad = 0)

plt.show()
figure = ax.get_figure()
figure.savefig('wordcloud.png', dpi=500)

```


```

[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data] /home/jovyan/nltk_data...
[nltk_data] Package averaged_perceptron_tagger is already up-to-
[nltk_data] date!


```




```
[12]: plt.figure(figsize = (10,6))
ax = sns.countplot(x = 'Class Name', data = test,palette = 'Paired')
plt.xticks(rotation = 45)
plt.show()
figure = ax.get_figure()
figure.savefig('category.png', dpi=500)
```




```
[13]: plt.figure(figsize = (10,6))
ax = sns.countplot(x = 'Department Name', data = test,palette = 'Spectral')
plt.xticks(rotation = 45)
plt.show()
figure = ax.get_figure()
figure.savefig('Department.png', dpi=500)
```



1.4.2 Unigram and Bigram


```
[14]: plt.figure(figsize = (12,10))
sns.barplot(data = df_freq[:10],y = 'freq',x = 'word',palette = 'Spectral').
    ↪set_title('Top 10 unigrams in review after data preprocessing')
plt.xticks(rotation = 45)
plt.savefig('Frequent Words.png', dpi=400)
```



```
[15]: from sklearn.feature_extraction.text import CountVectorizer
def top_n_ngram(corpus,n = None,ngram = 1):
    vec = CountVectorizer(stop_words = 'english',ngram_range=(ngram,ngram)).
    ↪fit(corpus)
    bag_of_words = vec.transform(corpus) #Have the count of all the words for
    ↪each review
    sum_words = bag_of_words.sum(axis =0) #Calculates the count of all the word
    ↪in the whole review
    words_freq = [(word,sum_words[0,idx]) for word,idx in vec.vocabulary_.
    ↪items()]
    words_freq = sorted(words_freq,key = lambda x:x[1],reverse = True)
    return words_freq[:n]

common_words = top_n_ngram(y, 10,2)
data = pd.DataFrame(common_words, columns = ['ReviewText' , 'count'])
plt.figure(figsize =(10,10))
bigram = data.groupby('ReviewText').sum()['count'].sort_values(ascending=False)
```

```
t = pd.DataFrame(data = {'words':bigram.index,'freq':bigram[:]})
sns.barplot(data = t,x = 'words',y = 'freq',palette = 'Spectral').set_title('Top
↳10 bigrams in review after data preprocessing')
plt.xticks(rotation = 45)
plt.savefig('Frequent Words bigram.jpg', dpi=400)
```




1.4.3 Polarity plot

```
[18]: from textblob import *
test['polarity'] = test['Review Text'].map(lambda text: TextBlob(text).
↳sentiment.polarity)
```

```
test['polarity']
```


```
[18]: 0      0.633333
      1      0.339583
      2      0.073675
      3      0.550000
      4      0.512891
      ...
     23481    0.552667
     23482    0.091667
     23483    0.414286
     23484    0.322222
     23485    0.413889
      Name: polarity, Length: 22641, dtype: float64
```

```
[19]: plt.figure(figsize = (8,10))
      ax = sns.histplot(test, x = 'polarity',hue="Rating")
      figure = ax.get_figure()
      figure.savefig('polarity.png', dpi=500)
```



```
[20]: df = test
negative = (len(df.loc[df.polarity < 0, ['Review Text']].values)/len(df))*100
positive = (len(df.loc[df.polarity > 0.5, ['Review Text']].values)/len(df))*100
neutral = len(df.loc[df.polarity > 0, ['Review Text']].values) - len(df.loc[df.
    ↳polarity > 0.5, ['Review Text']].values)
neutral = neutral/len(df)*100
plt.figure(figsize =(8, 8))
```

```
plt.pie([positive,negative,neutral], labels =_
↳['Positive','Negative','Neutral'],shadow = False, autopct='%2.2f%%',explode_
↳= (0, 0.1, 0))
plt.savefig('Percentage.jpg', dpi=500)
```



1.5 Text representation

1.5.1 3.1 One-hot encoding

```
[21]: # train_vec=np.concatenate([np.array(train_vec),df_train["Recommended IND"].
↳to_numpy().reshape(len(train_vec),1)],axis=1)
def one_hot_encoding(df):
```



```

X_train,X_test,y_train,y_test=train_test_split(df[["Review_
↪Text_processed"]],df["Rating"],test_size=0.2,random_state=123)
cv = CountVectorizer(binary=True,min_df=1,max_df=17448)
cv.fit(X_train["Review Text_processed"])
X_train_encoding=cv.transform(X_train["Review Text_processed"])
X_train_encoding=pd.DataFrame(X_train_encoding.A, columns=cv.
↪get_feature_names(),index=X_train.index)
X_test_encoding=cv.transform(X_test["Review Text_processed"])
X_test_encoding=pd.DataFrame(X_test_encoding.A, columns=cv.
↪get_feature_names(),index=X_test.index)

return X_train_encoding,X_test_encoding,y_train,y_test

```

```

[22]: X_train_oh,X_test_oh,y_train_oh,y_test_oh=one_hot_encoding(df_lf)
X_train_oh.head()

```

```

[22]:      aa  ab  abdomen  abil  abl  absolut  abstract  abt  abund  ac  ...  \
10365   0   0         0   0   0         0         0   0   0   0  ...
2579    0   0         0   0   0         0         0   0   0   0  ...
17002   0   0         0   0   0         0         0   0   0   0  ...
4121    0   0         0   0   0         0         0   0   0   0  ...
5440    0   0         0   0   0         0         0   0   0   0  ...

```

```

      yoke  young  younger  yr  yummi  zero  zip  zipper  zone  zoom
10365    0     0         0  0     0     0   0     0     0   0
2579     0     0         0  0     0     0   0     0     0   0
17002    0     0         0  0     0     0   0     0     0   0
4121     0     0         0  0     0     0   0     0     0   0
5440     0     0         0  0     0     0   0     0     0   0

```

[5 rows x 3360 columns]

1.5.2 3.2 TF-IDF

```

[23]: def tf_vectorization(df):
      X_train,X_test,y_train,y_test=train_test_split(df[["Review_
↪Text_processed"]],df["Rating"],test_size=0.2,random_state=123)
      tfidf = TfidfVectorizer(
          strip_accents=None,
          lowercase=True,
          preprocessor=None, # applied preprocessor in Data Cleaning
          tokenizer=None,
          use_idf=True,
          norm="l2",
          smooth_idf=True,
          min_df=2,
      ).fit(X_train["Review Text_processed"])

```

```

X_train_encoding=tfidf.transform(X_train["Review Text_processed"])
X_train_encoding=pd.DataFrame(X_train_encoding.A, columns=tfidf.
↳get_feature_names(),index=X_train.index)
X_test_encoding=tfidf.transform(X_test["Review Text_processed"])
X_test_encoding=pd.DataFrame(X_test_encoding.A, columns=tfidf.
↳get_feature_names(),index=X_test.index)

return X_train_encoding,X_test_encoding,y_train,y_test

```

```

[24]: X_train_tf,X_test_tf,y_train_tf,y_test_tf=tf_vectorization(df_lf)
X_test_tf.shape

```

```

[24]: (4529, 3355)

```

1.5.3 3.3 Word embedding

```

[25]: def word2vec_prep(tagged_data):
#     tagged_data=[word_tokenize(_d) for i, _d in
↳enumerate(df_train["Text_processed"])]
    word_counts=Counter(itertools.chain(*tagged_data)) #dict
#     vocabulary_inv=[x[0] for x in word_counts.most_common()]
#     vocabulary = {x: i for i, x in enumerate(vocabulary_inv)}
    return word_counts

def embedding_weights(vocabulary_inv,inp_data):
    model_name = "embedding"
    model_name = os.path.join(model_name)
    num_workers = 8 # Number of threads to run in parallel
    downsampling = 1e-3 # Downsample setting for frequent words
    sentences = [[vocabulary_inv[w] for w in s] for s in inp_data]
    size_features=200
    min_word_count=2
    context=5
    downsampling = 1e-3
    embedding_model = word2vec.Word2Vec(sentences, workers=num_workers,
                                       sg=0,
                                       size=size_features,
                                       min_count=min_word_count,
                                       window=context,
                                       sample=downsampling)
    embedding_model.init_sims(replace=True) #clean from RAM
    embedding_weights = np.zeros((len(vocabulary_inv), size_features))
    for i in range(len(vocabulary_inv)):
        word = vocabulary_inv[i]
        if word in embedding_model:
            embedding_weights[i] = embedding_model[word]

```

```

        else:
            embedding_weights[i] = np.random.uniform(-0.1, 0.1, size_features)

    return embedding_weights

def train_test_word2vec(tagged_train_data, tagged_test_data, weights, vocabulary):
    train_vec = []
    for doc in tagged_train_data:
        vec = 0
        for w in doc:
            vec += weights[vocabulary[w]]
        vec = vec / len(doc)
        train_vec.append(vec)

    test_vec = []
    for doc in tagged_test_data:
        vec = 0
        length = 0
        for w in doc:
            try:
                vec += weights[vocabulary[w]]
                length += 1
            except:
                continue
        vec = vec / length
        test_vec.append(vec)

    return np.array(train_vec), np.array(test_vec)

def word_embedding(df):
    df_train, df_test = train_test_split(df, test_size=0.2, shuffle=True)
    tagged_data = [word_tokenize(_d) for i, _d in enumerate(df_train["Review_
→Text_processed"])]
    word_counts = Counter(itertools.chain(*tagged_data))
    vocabulary_inv = [x[0] for x in word_counts.most_common()]
    vocabulary = {x: i for i, x in enumerate(vocabulary_inv)}
    inp_data = [[vocabulary[word] for word in text] for text in tagged_data]

    print("-----generating weights-----")
    weights = embedding_weights(vocabulary_inv, inp_data)

    tagged_train_data = [word_tokenize(_d) for i, _d in
→enumerate(df_train["Review Text_processed"])]
    tagged_test_data = [word_tokenize(_d) for i, _d in
→enumerate(df_test["Review Text_processed"])]

```

```

    print("-----generating train/test dataset-----")
    ↵
    ↪train_vec,test_vec=train_test_word2vec(tagged_train_data,tagged_test_data,weights,vocabular
#     train_vec=np.concatenate([train_vec,df_train["Recommended IND"].
    ↪to_numpy().reshape(len(train_vec),1)],axis=1)
#     test_vec=np.concatenate([test_vec,df_test["Recommended IND"].to_numpy().
    ↪reshape(len(test_vec),1)],axis=1)

    return train_vec,test_vec,df_train.Rating,df_test.Rating

```

```
[26]: X_train_we,X_test_we,y_train_we,y_test_we=word_embedding(df_lf)
```

```

-----generating weights-----
-----generating train/test dataset-----

```

1.6 Sentiment Analysis Model

1.6.1 BernoulliNB with Bag of words

```
[27]: from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import train_test_split
count = CountVectorizer()
X_count = count.fit_transform(reviews["Review Text_processed"]).toarray()
y=reviews['Recommended IND']

```

```
[28]: X_train, X_test, y_train, y_test = train_test_split(X_count, y, test_size = 0.
    ↪20, random_state = 0)
```

```
[29]: from sklearn.naive_bayes import BernoulliNB
classifier = BernoulliNB()
classifier.fit(X_train, y_train)
```

```
[29]: BernoulliNB(alpha=1.0, binarize=0.0, class_prior=None, fit_prior=True)
```

```
[30]: y_pred = classifier.predict(X_test)
y_train_pred = classifier.predict(X_train)
from sklearn.metrics import accuracy_score
from sklearn import metrics
acc1 = accuracy_score(y_test, y_pred)
score = accuracy_score(y_train,y_train_pred)
score
print("Accuracy of the classifier using Bag of words: ",acc1)
print("Confusion matrix is :\n",metrics.confusion_matrix(y_test,y_pred))
print("Classification report: \n" ,metrics.classification_report(y_test,y_pred))

```

Accuracy of the classifier using Bag of words: 0.8743652020313535
 Confusion matrix is :

```
[[ 625  252]
 [ 317 3335]]
Classification report:

```

	precision	recall	f1-score	support
0	0.66	0.71	0.69	877
1	0.93	0.91	0.92	3652
avg / total	0.88	0.87	0.88	4529


```
[31]: score
```

```
[31]: 0.8894655477031802
```

1.6.2 Logistic Regression with Bag of Words

```
[32]: res_train = []
res_test = []
c_range = np.logspace(-4,4,25)
import matplotlib.pyplot as plt
import numpy as np
from sklearn.linear_model import LogisticRegression
for cc in c_range:
    clf = LogisticRegression(random_state = 42, C=cc,
    ↪penalty='l2',solver='liblinear').fit(X_train, y_train)
    res_train.append(clf.score(X_train,y_train))
    res_test.append(clf.score(X_test,y_test))

plt.semilogx(c_range,res_train,label='train')
plt.semilogx(c_range,res_test,label='val')
plt.xlabel('c')
plt.ylabel('score')
plt.legend()
plt.show()
print(f'Best of c for the model is about ',[x for _,x in
    ↪sorted(zip(res_test,c_range))][-1])
```



Best of c for the model is about 0.46415888336127775

```
[33]: from sklearn.linear_model import LogisticRegression
      clf = LogisticRegression(random_state = 42, C=0.46,
      ↪penalty='l2',solver='liblinear').fit(X_train, y_train)
      print(clf.score(X_train,y_train))
      print(clf.score(X_test,y_test))
      y_pred = clf.predict(X_test)
      acc2 = accuracy_score(y_test, y_pred)
```

0.9237522084805654

0.8827555751821594

```
[34]: print("Accuracy of the LogisticRegression using Bag of words: ",acc2)
      print("Confusion matrix is :\n",metrics.confusion_matrix(y_test,y_pred))
      print("Classification report: \n" ,metrics.classification_report(y_test,y_pred))
```

Accuracy of the LogisticRegression using Bag of words: 0.8827555751821594

Confusion matrix is :

```
[[ 502  375]
```

```
 [ 156 3496]]
```

Classification report:

	precision	recall	f1-score	support
0	0.76	0.57	0.65	877


1	0.90	0.96	0.93	3652
avg / total	0.88	0.88	0.88	4529

1.6.3 Random Forest with Bag of Words

```
[35]: import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier
depth_list = np.arange(2,30,2)
acc_train, acc_test = [], []
for dp in depth_list:

    clf = RandomForestClassifier(max_depth=dp,random_state=0).
    →fit(X_train,y_train)
    acc_train.append(np.sum(clf.predict(X_train)==y_train)/len(X_train))
    acc_test.append(np.sum(clf.predict(X_test)==y_test)/len(X_test))

plt.plot(depth_list,acc_train)
plt.plot(depth_list,acc_test)
plt.show()
```



```
[36]: from sklearn.ensemble import RandomForestClassifier
clf = RandomForestClassifier(max_depth=26,random_state=0).fit(X_train,y_train)
print(np.sum(clf.predict(X_train)==y_train)/len(X_train))
print(np.sum(clf.predict(X_test)==y_test)/len(X_test))
```

```
y_pred = clf.predict(X_test)
acc3 = accuracy_score(y_test, y_pred)
```

```
0.8847725265017667
0.8271141532347096
```

```
[37]: print("Accuracy of the Random Forest using Bag of words: ",acc3)
      print("Confusion matrix is :\n",metrics.confusion_matrix(y_test,y_pred))
      print("Classification report: \n" ,metrics.classification_report(y_test,y_pred))
```

```
Accuracy of the Random Forest using Bag of words: 0.8271141532347096
```

```
Confusion matrix is :
```

```
[[ 115  762]
 [  21 3631]]
```

```
Classification report:
```

	precision	recall	f1-score	support
0	0.85	0.13	0.23	877
1	0.83	0.99	0.90	3652
avg / total	0.83	0.83	0.77	4529

1.6.4 Linear SVC with Bag of Words

```
[38]: ## Final method
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.svm import LinearSVC, SVC
      from sklearn.pipeline import make_pipeline
      from sklearn.preprocessing import StandardScaler, OneHotEncoder
      from sklearn.model_selection import train_test_split
      clf = make_pipeline(StandardScaler(),LinearSVC(random_state=0, tol=1e-4,C=0.
      ↪0001)).fit(X_train,y_train)

      #clf = LogisticRegression(max_iter=100000000).fit(train_cat_vec,
      ↪df_train["label"])

      print(np.sum(clf.predict(X_train)==y_train)/len(X_train))
      print(np.sum(clf.predict(X_test)==y_test)/len(X_test))
```

```
0.9260159010600707
0.870611614042835
```

```
[39]: clf = make_pipeline(StandardScaler(),LinearSVC(random_state=0, tol=1e-4,C=0.
      ↪00001)).fit(X_train,y_train)
```



```
#clf = LogisticRegression(max_iter=100000000).fit(train_cat_vec,
↳ df_train["label"])
```

```
print(np.sum(clf.predict(X_train)==y_train)/len(X_train))
print(np.sum(clf.predict(X_test)==y_test)/len(X_test))
```

0.9022195229681979
0.8703908147493928

```
[40]: y_pred = clf.predict(X_test)
acc4 = accuracy_score(y_test, y_pred)
print("Accuracy of the LinearSVC using Bag of words: ",acc4)
print("Confusion matrix is :\n",metrics.confusion_matrix(y_test,y_pred))
print("Classification report: \n" ,metrics.classification_report(y_test,y_pred))
```

Accuracy of the LinearSVC using Bag of words: 0.8703908147493928

Confusion matrix is :

```
[[ 635  242]
 [ 345 3307]]
```

Classification report:

	precision	recall	f1-score	support
0	0.65	0.72	0.68	877
1	0.93	0.91	0.92	3652
avg / total	0.88	0.87	0.87	4529

1.6.5 Multinomial NB with TF-IDF

```
[41]: from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
tfidf = TfidfVectorizer(max_features = 5000,ngram_range=(1, 2))
X_tfidf = tfidf.fit_transform(reviews["Review Text_processed"]).toarray()
```

```
[42]: from sklearn import feature_selection

y=reviews['Recommended IND']
X_names = tfidf.get_feature_names()
p_value_limit = 0.997 # 3 sigma cutoff
df_features = pd.DataFrame()

for cat in np.unique(y):
    chi2, p = feature_selection.chi2(X_tfidf, y==cat)
    df_features = df_features.append(pd.DataFrame(
        {"feature":X_names, "score":1-p, "label":cat}))
    df_features = df_features.sort_values(["label","score"],
        ascending=[True,False])
```

```
df_features = df_features[df_features["score"]>p_value_limit]

X_names = df_features["feature"].unique().tolist()
print(len(X_names))
```

238

```
[43]: X_train, X_test, y_train, y_test = train_test_split(X_tfidf, y, test_size = 0.
      ↪20, random_state = 0)
```

```
[44]: X_train
```

```
[44]: array([[0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            ...,
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.]])
```

```
[45]: classifier = MultinomialNB()
      classifier.fit(X_train, y_train)
```

```
[45]: MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)
```

```
[46]: y_pred = classifier.predict(X_test)
      y_train_pred = classifier.predict(X_train)
      acc5 = accuracy_score(y_test, y_pred)
      score = accuracy_score(y_train,y_train_pred)
```

```
[47]: score
```

```
[47]: 0.8849381625441696
```

```
[48]: print("Accuracy of the classifier using TF-IDF: ",acc5)
      print("Confusion matrix is :\n",metrics.confusion_matrix(y_test,y_pred))
      print("Classification report: \n" ,metrics.classification_report(y_test,y_pred))
```

Accuracy of the classifier using TF-IDF: 0.8593508500772797

Confusion matrix is :

```
[[ 289  588]
 [  49 3603]]
```

Classification report:

	precision	recall	f1-score	support
0	0.86	0.33	0.48	877
1	0.86	0.99	0.92	3652


avg / total	0.86	0.86	0.83	4529
-------------	------	------	------	------

1.6.6 Logistic Regression with TF-IDF

```
[49]: res_train = []
res_test = []
c_range = np.logspace(-4,4,25)
import matplotlib.pyplot as plt
import numpy as np

for cc in c_range:
    clf = LogisticRegression(random_state = 42, C=cc,
    ↪penalty='l2',solver='liblinear').fit(X_train, y_train)
    res_train.append(clf.score(X_train,y_train))
    res_test.append(clf.score(X_test,y_test))

plt.semilogx(c_range,res_train,label='train')
plt.semilogx(c_range,res_test,label='val')
plt.xlabel('c')
plt.ylabel('score')
plt.legend()
plt.show()
print(f'Best of c for the model is about ',[x for _,x in
    ↪sorted(zip(res_test,c_range))][-1])
```



Best of c for the model is about 4.641588833612772

```
[52]: clf = LogisticRegression(random_state = 42, C=4.64,  
    ↪penalty='l2',solver='liblinear').fit(X_train, y_train)  
print(clf.score(X_train,y_train))  
print(clf.score(X_test,y_test))  
y_pred = clf.predict(X_test)  
acc6 = accuracy_score(y_test, y_pred)
```

0.9408679328621908

0.8876131596378891

```
[53]: print("Accuracy of the LogisticRegression using TF-IDF: ",acc6)  
print("Confusion matrix is :\n",metrics.confusion_matrix(y_test,y_pred))  
print("Classification report: \n" ,metrics.classification_report(y_test,y_pred))
```

Accuracy of the LogisticRegression using TF-IDF: 0.8876131596378891

Confusion matrix is :

[[515 362]


[147 3505]]

Classification report:

	precision	recall	f1-score	support
0	0.78	0.59	0.67	877
1	0.91	0.96	0.93	3652
avg / total	0.88	0.89	0.88	4529

1.6.7 Random Forest with TF-IDF

```
[54]: import matplotlib.pyplot as plt  
depth_list = np.arange(2,20,2)  
acc_train, acc_test = [], []  
for dp in depth_list:  
  
    clf = RandomForestClassifier(max_depth=dp,random_state=0).  
    ↪fit(X_train,y_train)  
    acc_train.append(np.sum(clf.predict(X_train)==y_train)/len(X_train))  
    acc_test.append(np.sum(clf.predict(X_test)==y_test)/len(X_test))  
  
plt.plot(depth_list,acc_train)  
plt.plot(depth_list,acc_test)  
plt.show()
```



```
[55]: clf = RandomForestClassifier(max_depth=20,random_state=0).fit(X_train,y_train)
print(clf.score(X_train,y_train))
print(clf.score(X_test,y_test))
y_pred = clf.predict(X_test)
acc7 = accuracy_score(y_test, y_pred)
print("Accuracy of the Random Forest using TF-IDF: ",acc7)
print("Confusion matrix is :\n",metrics.confusion_matrix(y_test,y_pred))
print("Classification report: \n" ,metrics.classification_report(y_test,y_pred))
```

0.8646201413427562

0.8222565687789799

Accuracy of the Random Forest using TF-IDF: 0.8222565687789799

Confusion matrix is :

```
[[ 94 783]
```

```
[ 22 3630]]
```

Classification report:

	precision	recall	f1-score	support
0	0.81	0.11	0.19	877
1	0.82	0.99	0.90	3652
avg / total	0.82	0.82	0.76	4529

1.6.8 Linear SVC with TF-IDF

```
[56]: ## Final method
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import LinearSVC, SVC
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.model_selection import train_test_split
clf = make_pipeline(StandardScaler(),LinearSVC(random_state=0, tol=1e-4,C=0.
    ↳0001)).fit(X_train,y_train)

#clf = LogisticRegression(max_iter=100000000).fit(train_cat_vec,
    ↳df_train["label"])

print(np.sum(clf.predict(X_train)==y_train)/len(X_train))
print(np.sum(clf.predict(X_test)==y_test)/len(X_test))
```

0.9479902826855123

0.8827555751821594

```
[57]: y_pred = clf.predict(X_test)
acc8 = accuracy_score(y_test, y_pred)
print("Accuracy of the LinearSVC using TF-IDF: ",acc8)
print("Confusion matrix is :\n",metrics.confusion_matrix(y_test,y_pred))
print("Classification report: \n" ,metrics.classification_report(y_test,y_pred))
```

Accuracy of the LinearSVC using TF-IDF: 0.8827555751821594

Confusion matrix is :

```
[[ 580  297]
```

```
 [ 234 3418]]
```

Classification report:

	precision	recall	f1-score	support
0	0.71	0.66	0.69	877
1	0.92	0.94	0.93	3652
avg / total	0.88	0.88	0.88	4529

1.6.9 word2vec

```
[58]: import multiprocessing

from gensim.models import Word2Vec

w2v_model = Word2Vec()
w2v_model.build_vocab(reviews["Review Text_processed"], progress_per=10000)
```

```
w2v_model.train(reviews["Review Text_processed"], total_examples=w2v_model.
↳corpus_count, epochs=30, report_delay=1)

w2v_model.init_sims(replace=True)
```

```
[59]: import numpy as np
from numpy import dot
from numpy.linalg import norm
def document_vector_w2v(doc):
    """Create document vectors by averaging word vectors. Remove
↳out-of-vocabulary words."""
    doc = [word for word in doc if word in w2v_model.wv.vocab]
    return np.mean(w2v_model[doc], axis=0)
```

```
[60]: X_w2v=reviews["Review Text_processed"].apply(lambda x: ' '.join(x)).
↳apply(document_vector_w2v)
y_w2v = reviews['Recommended IND']
```

```
[61]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_w2v, y_w2v, test_size = 0.
↳20, random_state = 0)
```


```
[62]: X_train, X_test, y_train, y_test = list(X_train), list(X_test), list(y_train),
↳list(y_test)
```

1.6.10 Logistic Regression with word2vec

```
[63]: res_train = []
res_test = []
c_range = np.logspace(-4,4,25)
import matplotlib.pyplot as plt
import numpy as np
from sklearn.linear_model import LogisticRegression
for cc in c_range:
    clf = LogisticRegression(random_state = 42, C=cc,
↳penalty='l2',solver='liblinear').fit(X_train, y_train)
    res_train.append(clf.score(X_train,y_train))
    res_test.append(clf.score(X_test,y_test))

plt.semilogx(c_range,res_train,label='train')
plt.semilogx(c_range,res_test,label='val')
plt.xlabel('c')
plt.ylabel('score')
plt.legend()
plt.show()
```

```
print(f'Best of c for the model is about ', [x for _, x in
↪sorted(zip(res_test, c_range))][-1])
```



Best of c for the model is about 4641.588833612773

```
[66]: from sklearn.metrics import accuracy_score
from sklearn import metrics
clf = LogisticRegression(random_state = 42, C=1000,
↪penalty='l2', solver='liblinear').fit(X_train, y_train)
print(clf.score(X_train, y_train))
print(clf.score(X_test, y_test))
y_pred = clf.predict(X_test)
acc9 = accuracy_score(y_test, y_pred)
```

0.8219412544169611
0.8092294104658865

```
[67]: print("Accuracy of the LogisticRegression using word2vec: ", acc9)
print("Confusion matrix is :\n", metrics.confusion_matrix(y_test, y_pred))
print("Classification report: \n", metrics.classification_report(y_test, y_pred))
```

Accuracy of the LogisticRegression using word2vec: 0.8092294104658865
Confusion matrix is :
[[29 848]
[16 3636]]

Classification report:


	precision	recall	f1-score	support
0	0.64	0.03	0.06	877
1	0.81	1.00	0.89	3652
avg / total	0.78	0.81	0.73	4529

1.6.11 Random Forest with word2vec

```
[21]: from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import LinearSVC, SVC
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
import matplotlib.pyplot as plt
depth_list = np.arange(2,30,2)
acc_train, acc_test = [], []
for dp in depth_list:

    clf = RandomForestClassifier(max_depth=dp,random_state=0).
    ↪fit(X_train,y_train)
    acc_train.append(np.sum(clf.predict(X_train)==y_train)/len(X_train))
    acc_test.append(np.sum(clf.predict(X_test)==y_test)/len(X_test))

plt.plot(depth_list,acc_train)
plt.plot(depth_list,acc_test)
plt.show()
```



```
[68]: clf = RandomForestClassifier(max_depth=20,random_state=0).fit(X_train,y_train)
print(clf.score(X_train,y_train))
print(clf.score(X_test,y_test))
y_pred = clf.predict(X_test)
acc10 = accuracy_score(y_test, y_pred)
print("Accuracy of the Random Forest using word2vec: ",acc10)
print("Confusion matrix is :\n",metrics.confusion_matrix(y_test,y_pred))
print("Classification report: \n" ,metrics.classification_report(y_test,y_pred))
```

0.969357332155477

0.7895782733495252

Accuracy of the Random Forest using word2vec: 0.7895782733495252

Confusion matrix is :

```
[[ 68 809]
```

```
[ 144 3508]]
```

Classification report:

	precision	recall	f1-score	support
0	0.32	0.08	0.12	877
1	0.81	0.96	0.88	3652
avg / total	0.72	0.79	0.73	4529

1.6.12 Linear SVC with word2vec

```
[69]: ## Final method
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import LinearSVC, SVC
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.model_selection import train_test_split
clf = make_pipeline(StandardScaler(),LinearSVC(random_state=0, tol=1e-4,C=0.
    ↳00001)).fit(X_train,y_train)

#clf = LogisticRegression(max_iter=100000000).fit(train_cat_vec,
    ↳df_train["label"])

print(np.sum(clf.predict(X_train)==y_train)/len(X_train))
print(np.sum(clf.predict(X_test)==y_test)/len(X_test))
```

0.8197327738515902

0.8061382203576949

```
[70]: y_pred = clf.predict(X_test)
acc11 = accuracy_score(y_test, y_pred)
print("Accuracy of the LinearSVC using word2vec: ",acc11)
print("Confusion matrix is :\n",metrics.confusion_matrix(y_test,y_pred))
print("Classification report: \n" ,metrics.classification_report(y_test,y_pred))
```

Accuracy of the LinearSVC using word2vec: 0.8061382203576949

Confusion matrix is :

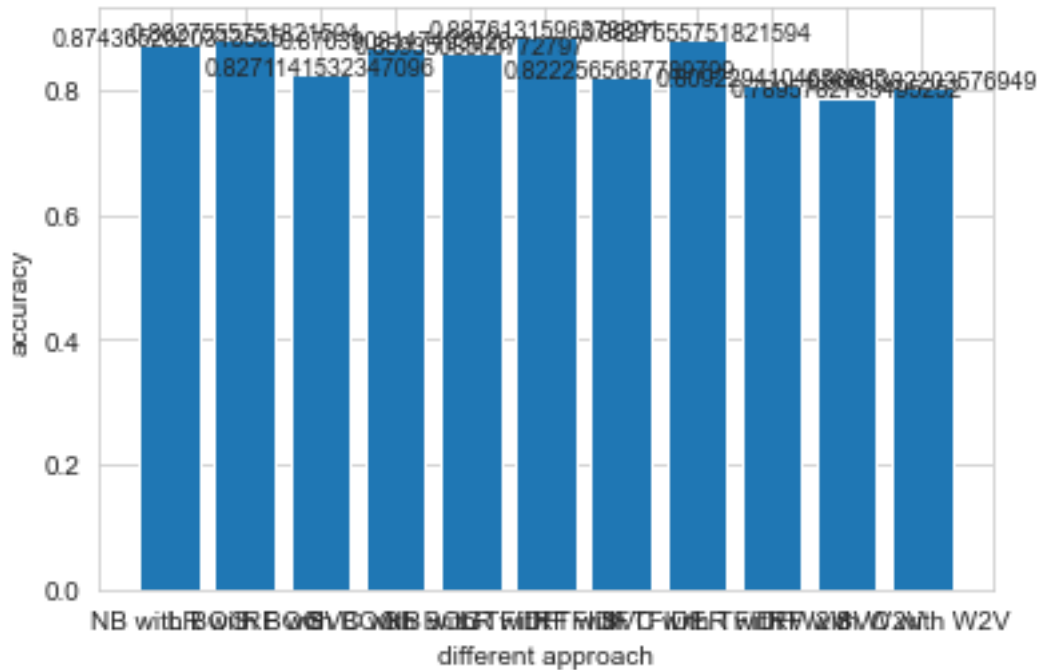
```
[[ 47 830]
```

```
[ 48 3604]]
```

Classification report:

	precision	recall	f1-score	support
0	0.49	0.05	0.10	877
1	0.81	0.99	0.89	3652
avg / total	0.75	0.81	0.74	4529

```
[71]: import matplotlib.pyplot as plt
index=['NB with BOG','LR with BOG','RF with BOG','SVC with BOG','NB with
↳TFIDF','LR with TFIDF','RF with TFIDF','SVC with TFIDF','LR with W2V','RF
↳with W2V','SVC with W2V']
acc = [acc1,acc2,acc3,acc4,acc5,acc6,acc7,acc8,acc9,acc10,acc11]
plt.bar(index,acc)
for index,value in enumerate(acc):
    plt.text(index,value, str(value),fontsize=9,horizontalalignment='center')
plt.ylabel('accuracy')
plt.xlabel('different approach')
plt.show()
```



```
[72]: index=['NB with BOG','LR with BOG','RF with BOG','SVC with BOG','NB with
↳TFIDF','LR with TFIDF','RF with TFIDF','SVC with TFIDF','LR with W2V','RF
↳with W2V','SVC with W2V']
acc = [acc1,acc2,acc3,acc4,acc5,acc6,acc7,acc8,acc9,acc10,acc11]
d = {'approach':index,'accuracy':acc}
df=pd.DataFrame(data=d)
df
```

```
[72]:      approach  accuracy
0      NB with BOG  0.874365
1      LR with BOG  0.882756
2      RF with BOG  0.827114
3      SVC with BOG  0.870391
4      NB with TFIDF  0.859351
5      LR with TFIDF  0.887613
6      RF with TFIDF  0.822257
7      SVC with TFIDF  0.882756
8      LR with W2V   0.809229
9      RF with W2V   0.789578
10     SVC with W2V   0.806138
```

1.6.13 Deep Learning

```
[73]: import tensorflow as tf
      from tensorflow.keras.preprocessing.text import Tokenizer
      from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
[74]: tokenizer = Tokenizer(num_words = 3000)
      tokenizer.fit_on_texts(reviews["Review Text_processed"])
```

```
[75]: sequences = tokenizer.texts_to_sequences(reviews["Review Text_processed"])
      padded = pad_sequences(sequences, padding='post')
```

```
[76]: word_index = tokenizer.word_index
      count = 0
      for i,j in word_index.items():
          if count == 11:
              break
          print(i,j)
          count = count+1
```

```
dress 1
fit 2
love 3
size 4
look 5
top 6
wear 7
like 8
color 9
would 10
great 11
```

```
[77]: embedding_dim = 64
      model = tf.keras.Sequential([
          tf.keras.layers.Embedding(3000, embedding_dim),
          tf.keras.layers.GlobalAveragePooling1D(),
          tf.keras.layers.Dense(6, activation='relu'),
          tf.keras.layers.Dense(1, activation='sigmoid')
      ])

      model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, None, 64)	192000

```

global_average_pooling1d (Gl (None, 64)                0
-----
dense (Dense)                (None, 6)                390
-----
dense_1 (Dense)              (None, 1)                7
=====
Total params: 192,397
Trainable params: 192,397
Non-trainable params: 0
-----

```

```

[78]: num_epochs = 10

model.compile(loss='binary_crossentropy',optimizer='adam',metrics=['accuracy'])

```

```

[79]: model.fit(padded,y,epochs= num_epochs)

```

```

Epoch 1/10
708/708 [=====] - 3s 4ms/step - loss: 0.4099 -
accuracy: 0.8247
Epoch 2/10
708/708 [=====] - 3s 4ms/step - loss: 0.2554 -
accuracy: 0.8906
Epoch 3/10
708/708 [=====] - 3s 4ms/step - loss: 0.2294 -
accuracy: 0.9037
Epoch 4/10
708/708 [=====] - 3s 4ms/step - loss: 0.2166 -
accuracy: 0.9103
Epoch 5/10
708/708 [=====] - 3s 4ms/step - loss: 0.2079 -
accuracy: 0.9146
Epoch 6/10
708/708 [=====] - 3s 4ms/step - loss: 0.2000 -
accuracy: 0.9189
Epoch 7/10
708/708 [=====] - 3s 4ms/step - loss: 0.1934 -
accuracy: 0.9216
Epoch 8/10
708/708 [=====] - 3s 4ms/step - loss: 0.1875 -
accuracy: 0.9231
Epoch 9/10
708/708 [=====] - 3s 4ms/step - loss: 0.1810 -
accuracy: 0.9262
Epoch 10/10
708/708 [=====] - 3s 4ms/step - loss: 0.1754 -
accuracy: 0.9275

```

```
[79]: <tensorflow.python.keras.callbacks.History at 0x266528c67b8>
```

```
[80]: sample_string = "I Will tell my friends for sure"
sample = tokenizer.texts_to_sequences(sample_string)
padded_sample = pad_sequences(sample, padding='post')
```

```
[81]: padded_sample.T
```

```
[81]: array([], shape=(0, 31), dtype=int32)
```

1.7 Rating Prediction Model

1.7.1 Neural Network (Keras)

```
[15]: def NN_fit(X_train,y_train):
    model=Sequential()
    n_cols=X_train.shape[1]
    model.add(Dense(11,activation='relu',input_shape=(n_cols,)))
    model.add(Dense(11,activation='relu'))
    # model.add(Dense(20,activation='relu'))
    # model.add(Dense(20,activation='relu'))

    model.add(Dense(5,activation='softmax'))
    model.
    ↪ compile(optimizer='adam',loss='categorical_crossentropy',metrics=["accuracy"])
    early_stopping_monitor=EarlyStopping(patience=3)
    model.fit(X_train,pd.get_dummies(y_train).to_numpy(),validation_split=0.
    ↪ 2,epochs=100,callbacks=[early_stopping_monitor])

    return model
```

```
[59]: model_test=NN_fit(X_train_we,y_train_we)
```

Epoch 1/100

453/453 [=====] - 1s 1ms/step - loss: 1.3091 - accuracy: 0.5473 - val_loss: 1.0176 - val_accuracy: 0.5904

Epoch 2/100

453/453 [=====] - 0s 870us/step - loss: 1.0120 - accuracy: 0.5975 - val_loss: 0.9920 - val_accuracy: 0.5945

Epoch 3/100

453/453 [=====] - 0s 944us/step - loss: 1.0046 - accuracy: 0.5898 - val_loss: 0.9819 - val_accuracy: 0.6050

Epoch 4/100

453/453 [=====] - 0s 915us/step - loss: 0.9809 - accuracy: 0.5974 - val_loss: 0.9740 - val_accuracy: 0.6064

Epoch 5/100

453/453 [=====] - 0s 923us/step - loss: 0.9772 -

accuracy: 0.6048 - val_loss: 0.9687 - val_accuracy: 0.6072
Epoch 6/100
453/453 [=====] - 0s 934us/step - loss: 0.9741 -
accuracy: 0.6042 - val_loss: 0.9667 - val_accuracy: 0.6089
Epoch 7/100
453/453 [=====] - 0s 883us/step - loss: 0.9861 -
accuracy: 0.6005 - val_loss: 0.9609 - val_accuracy: 0.6122
Epoch 8/100
453/453 [=====] - 0s 903us/step - loss: 0.9791 -
accuracy: 0.5987 - val_loss: 0.9574 - val_accuracy: 0.6130
Epoch 9/100
453/453 [=====] - 0s 887us/step - loss: 0.9759 -
accuracy: 0.6022 - val_loss: 0.9549 - val_accuracy: 0.6147
Epoch 10/100
453/453 [=====] - 0s 956us/step - loss: 0.9605 -
accuracy: 0.6056 - val_loss: 0.9525 - val_accuracy: 0.6125
Epoch 11/100
453/453 [=====] - 0s 933us/step - loss: 0.9662 -
accuracy: 0.6000 - val_loss: 0.9530 - val_accuracy: 0.6133
Epoch 12/100
453/453 [=====] - 0s 850us/step - loss: 0.9697 -
accuracy: 0.6053 - val_loss: 0.9485 - val_accuracy: 0.6144
Epoch 13/100
453/453 [=====] - 0s 885us/step - loss: 0.9723 -
accuracy: 0.5970 - val_loss: 0.9493 - val_accuracy: 0.6128
Epoch 14/100
453/453 [=====] - 0s 909us/step - loss: 0.9527 -
accuracy: 0.6079 - val_loss: 0.9453 - val_accuracy: 0.6169
Epoch 15/100
453/453 [=====] - 0s 923us/step - loss: 0.9569 -
accuracy: 0.6054 - val_loss: 0.9429 - val_accuracy: 0.6163
Epoch 16/100
453/453 [=====] - 0s 887us/step - loss: 0.9580 -
accuracy: 0.6056 - val_loss: 0.9433 - val_accuracy: 0.6163
Epoch 17/100
453/453 [=====] - 0s 915us/step - loss: 0.9529 -
accuracy: 0.6040 - val_loss: 0.9458 - val_accuracy: 0.6180
Epoch 18/100
453/453 [=====] - 0s 921us/step - loss: 0.9653 -
accuracy: 0.6051 - val_loss: 0.9418 - val_accuracy: 0.6199
Epoch 19/100
453/453 [=====] - 0s 967us/step - loss: 0.9476 -
accuracy: 0.6142 - val_loss: 0.9428 - val_accuracy: 0.6166
Epoch 20/100
453/453 [=====] - 0s 938us/step - loss: 0.9496 -
accuracy: 0.6104 - val_loss: 0.9414 - val_accuracy: 0.6221
Epoch 21/100
453/453 [=====] - 0s 986us/step - loss: 0.9403 -


```

accuracy: 0.6102 - val_loss: 0.9446 - val_accuracy: 0.6191
Epoch 22/100
453/453 [=====] - 0s 909us/step - loss: 0.9401 -
accuracy: 0.6124 - val_loss: 0.9392 - val_accuracy: 0.6197
Epoch 23/100
453/453 [=====] - 0s 959us/step - loss: 0.9544 -
accuracy: 0.6073 - val_loss: 0.9423 - val_accuracy: 0.6172
Epoch 24/100
453/453 [=====] - 0s 942us/step - loss: 0.9445 -
accuracy: 0.6115 - val_loss: 0.9380 - val_accuracy: 0.6213
Epoch 25/100
453/453 [=====] - 0s 964us/step - loss: 0.9454 -
accuracy: 0.6114 - val_loss: 0.9362 - val_accuracy: 0.6208
Epoch 26/100
453/453 [=====] - 0s 995us/step - loss: 0.9402 -
accuracy: 0.6148 - val_loss: 0.9390 - val_accuracy: 0.6188
Epoch 27/100
453/453 [=====] - 0s 931us/step - loss: 0.9390 -
accuracy: 0.6112 - val_loss: 0.9394 - val_accuracy: 0.6219
Epoch 28/100
453/453 [=====] - 0s 964us/step - loss: 0.9461 -
accuracy: 0.6182 - val_loss: 0.9363 - val_accuracy: 0.6230

```

```
[63]: np.unique(model_test.predict_classes(X_train_we))
```

```
[63]: array([0, 1, 2, 3, 4])
```

```

[77]: encoding_name=["One-hot+NN","TF-IDF+NN","Word_embedding+NN"]
data_pool=[[X_train_oh.to_numpy(),X_test_oh.to_numpy(),y_train_oh,y_test_oh],
            [X_train_tf.to_numpy(),X_test_tf.to_numpy(),y_train_tf,y_test_tf],
            [X_train_we,X_test_we,y_train_we,y_test_we]]

score_list=[]

for data in data_pool:
    clf=NN_fit(data[0],data[2])
    y_train_pred_nn=clf.predict_classes(data[0])+1
    y_test_pred_nn=clf.predict_classes(data[1])+1
    # train_score=clf.predict(data[0],pd.get_dummies(data[2]))[1]
    # test_score=clf.predict(data[1],pd.get_dummies(data[3]))[1]
    score_list.append([metrics.accuracy_score(data[2], y_train_pred_nn),
                       metrics.accuracy_score(data[3], y_test_pred_nn),
                       f1_score(data[2], y_train_pred_nn, average="macro"),
                       f1_score(data[3], y_test_pred_nn, average="macro")])

result_nn=pd.
    ↳DataFrame(score_list,columns=["Train_acc","Test_acc","Train_f1","Test_f1"],index=encoding_n

```

```
result_nn
```

```
Epoch 1/100
453/453 [=====] - 1s 2ms/step - loss: 1.1646 -
accuracy: 0.5546 - val_loss: 0.9048 - val_accuracy: 0.6335
Epoch 2/100
453/453 [=====] - 1s 1ms/step - loss: 0.8308 -
accuracy: 0.6567 - val_loss: 0.8895 - val_accuracy: 0.6332
Epoch 3/100
453/453 [=====] - 1s 1ms/step - loss: 0.7435 -
accuracy: 0.6954 - val_loss: 0.9008 - val_accuracy: 0.6296
Epoch 4/100
453/453 [=====] - 1s 1ms/step - loss: 0.6621 -
accuracy: 0.7345 - val_loss: 0.9335 - val_accuracy: 0.6174
Epoch 5/100
453/453 [=====] - 1s 1ms/step - loss: 0.6154 -
accuracy: 0.7598 - val_loss: 0.9643 - val_accuracy: 0.6119

/Users/yinuochen/opt/anaconda3/lib/python3.8/site-
packages/tensorflow/python/keras/engine/sequential.py:450: UserWarning:
`model.predict_classes()` is deprecated and will be removed after 2021-01-01.
Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your model
does multi-class classification (e.g. if it uses a `softmax` last-layer
activation). * `(model.predict(x) > 0.5).astype("int32")`, if your model does
binary classification (e.g. if it uses a `sigmoid` last-layer activation).
  warnings.warn("`model.predict_classes()` is deprecated and '

Epoch 1/100
453/453 [=====] - 1s 1ms/step - loss: 1.3463 -
accuracy: 0.5313 - val_loss: 0.9450 - val_accuracy: 0.6114
Epoch 2/100
453/453 [=====] - 1s 1ms/step - loss: 0.8831 -
accuracy: 0.6314 - val_loss: 0.8740 - val_accuracy: 0.6351
Epoch 3/100
453/453 [=====] - 1s 1ms/step - loss: 0.7815 -
accuracy: 0.6789 - val_loss: 0.8686 - val_accuracy: 0.6390
Epoch 4/100
453/453 [=====] - 1s 1ms/step - loss: 0.7146 -
accuracy: 0.7092 - val_loss: 0.8846 - val_accuracy: 0.6329
Epoch 5/100
453/453 [=====] - 1s 1ms/step - loss: 0.6659 -
accuracy: 0.7306 - val_loss: 0.9045 - val_accuracy: 0.6323
Epoch 6/100
453/453 [=====] - 1s 1ms/step - loss: 0.6409 -
accuracy: 0.7419 - val_loss: 0.9310 - val_accuracy: 0.6147
Epoch 1/100
453/453 [=====] - 1s 1ms/step - loss: 1.3001 -
accuracy: 0.5276 - val_loss: 1.0313 - val_accuracy: 0.5702
```

Epoch 2/100
453/453 [=====] - 0s 864us/step - loss: 1.0341 - accuracy: 0.5833 - val_loss: 1.0002 - val_accuracy: 0.6001

Epoch 3/100
453/453 [=====] - 0s 925us/step - loss: 0.9960 - accuracy: 0.6013 - val_loss: 0.9832 - val_accuracy: 0.6034

Epoch 4/100
453/453 [=====] - 0s 901us/step - loss: 0.9918 - accuracy: 0.5989 - val_loss: 0.9755 - val_accuracy: 0.6070

Epoch 5/100
453/453 [=====] - 0s 883us/step - loss: 0.9829 - accuracy: 0.5938 - val_loss: 0.9703 - val_accuracy: 0.6083

Epoch 6/100
453/453 [=====] - 0s 916us/step - loss: 0.9925 - accuracy: 0.5973 - val_loss: 0.9704 - val_accuracy: 0.6025

Epoch 7/100
453/453 [=====] - 0s 963us/step - loss: 0.9932 - accuracy: 0.5900 - val_loss: 0.9641 - val_accuracy: 0.6072

Epoch 8/100
453/453 [=====] - 0s 913us/step - loss: 0.9603 - accuracy: 0.6080 - val_loss: 0.9616 - val_accuracy: 0.6072

Epoch 9/100
453/453 [=====] - 0s 884us/step - loss: 0.9677 - accuracy: 0.6052 - val_loss: 0.9584 - val_accuracy: 0.6094

Epoch 10/100
453/453 [=====] - 0s 939us/step - loss: 0.9618 - accuracy: 0.6096 - val_loss: 0.9569 - val_accuracy: 0.6097

Epoch 11/100
453/453 [=====] - 0s 966us/step - loss: 0.9631 - accuracy: 0.6071 - val_loss: 0.9546 - val_accuracy: 0.6116

Epoch 12/100
453/453 [=====] - 0s 865us/step - loss: 0.9653 - accuracy: 0.6026 - val_loss: 0.9531 - val_accuracy: 0.6128

Epoch 13/100
453/453 [=====] - 0s 921us/step - loss: 0.9698 - accuracy: 0.6040 - val_loss: 0.9666 - val_accuracy: 0.6036

Epoch 14/100
453/453 [=====] - 0s 862us/step - loss: 0.9641 - accuracy: 0.6030 - val_loss: 0.9518 - val_accuracy: 0.6128

Epoch 15/100
453/453 [=====] - 0s 947us/step - loss: 0.9581 - accuracy: 0.6082 - val_loss: 0.9494 - val_accuracy: 0.6083

Epoch 16/100
453/453 [=====] - 0s 874us/step - loss: 0.9553 - accuracy: 0.6048 - val_loss: 0.9514 - val_accuracy: 0.6144

Epoch 17/100
453/453 [=====] - 0s 943us/step - loss: 0.9553 - accuracy: 0.6063 - val_loss: 0.9472 - val_accuracy: 0.6119

Epoch 18/100
453/453 [=====] - 0s 870us/step - loss: 0.9586 -
accuracy: 0.6084 - val_loss: 0.9459 - val_accuracy: 0.6163
Epoch 19/100
453/453 [=====] - 0s 875us/step - loss: 0.9546 -
accuracy: 0.6057 - val_loss: 0.9585 - val_accuracy: 0.6047
Epoch 20/100
453/453 [=====] - 0s 922us/step - loss: 0.9587 -
accuracy: 0.6069 - val_loss: 0.9447 - val_accuracy: 0.6125
Epoch 21/100
453/453 [=====] - 0s 985us/step - loss: 0.9448 -
accuracy: 0.6120 - val_loss: 0.9430 - val_accuracy: 0.6183
Epoch 22/100
453/453 [=====] - 0s 886us/step - loss: 0.9625 -
accuracy: 0.6104 - val_loss: 0.9439 - val_accuracy: 0.6161
Epoch 23/100
453/453 [=====] - 0s 895us/step - loss: 0.9440 -
accuracy: 0.6140 - val_loss: 0.9435 - val_accuracy: 0.6180
Epoch 24/100
453/453 [=====] - 0s 869us/step - loss: 0.9398 -
accuracy: 0.6147 - val_loss: 0.9412 - val_accuracy: 0.6174
Epoch 25/100
453/453 [=====] - 0s 931us/step - loss: 0.9396 -
accuracy: 0.6132 - val_loss: 0.9422 - val_accuracy: 0.6128
Epoch 26/100
453/453 [=====] - 0s 944us/step - loss: 0.9450 -
accuracy: 0.6061 - val_loss: 0.9404 - val_accuracy: 0.6205
Epoch 27/100
453/453 [=====] - 0s 912us/step - loss: 0.9483 -
accuracy: 0.6074 - val_loss: 0.9448 - val_accuracy: 0.6161
Epoch 28/100
453/453 [=====] - 0s 964us/step - loss: 0.9261 -
accuracy: 0.6167 - val_loss: 0.9389 - val_accuracy: 0.6177
Epoch 29/100
453/453 [=====] - 0s 869us/step - loss: 0.9386 -
accuracy: 0.6158 - val_loss: 0.9407 - val_accuracy: 0.6177
Epoch 30/100
453/453 [=====] - 0s 966us/step - loss: 0.9529 -
accuracy: 0.5990 - val_loss: 0.9439 - val_accuracy: 0.6243
Epoch 31/100
453/453 [=====] - 0s 924us/step - loss: 0.9330 -
accuracy: 0.6116 - val_loss: 0.9387 - val_accuracy: 0.6199
Epoch 32/100
453/453 [=====] - 0s 968us/step - loss: 0.9368 -
accuracy: 0.6113 - val_loss: 0.9370 - val_accuracy: 0.6232
Epoch 33/100
453/453 [=====] - 0s 871us/step - loss: 0.9373 -
accuracy: 0.6120 - val_loss: 0.9358 - val_accuracy: 0.6210

```
Epoch 34/100
453/453 [=====] - 0s 909us/step - loss: 0.9310 -
accuracy: 0.6158 - val_loss: 0.9379 - val_accuracy: 0.6241
Epoch 35/100
453/453 [=====] - 0s 939us/step - loss: 0.9329 -
accuracy: 0.6132 - val_loss: 0.9354 - val_accuracy: 0.6252
Epoch 36/100
453/453 [=====] - 0s 878us/step - loss: 0.9441 -
accuracy: 0.6079 - val_loss: 0.9361 - val_accuracy: 0.6232
Epoch 37/100
453/453 [=====] - 0s 875us/step - loss: 0.9368 -
accuracy: 0.6151 - val_loss: 0.9429 - val_accuracy: 0.6133
Epoch 38/100
453/453 [=====] - 0s 976us/step - loss: 0.9395 -
accuracy: 0.6095 - val_loss: 0.9362 - val_accuracy: 0.6172
```

```
[77]:
```

	Train_acc	Test_acc	Train_f1	Test_f1
One-hot+NN	0.755411	0.618238	0.636620	0.443728
TF-IDF+NN	0.741884	0.612497	0.621029	0.421230
Word_embedding+NN	0.613626	0.608302	0.356031	0.346883

1.7.2 LinearSVC

```
[88]: # from sklearn import svm
# from sklearn import metrics

# clf = svm.SVC(kernel='linear')
# clf.fit(X_train_oh,y_train_oh)
# y_pred = clf.predict(X_test_oh)
# print("Accuracy:",metrics.accuracy_score(y_test_oh, y_pred))
```

```
[50]: start_time = time.time()
# print(start_time)
clf = OneVsRestClassifier(SVC(kernel='linear', probability=False), n_jobs=-1)
clf.fit(X_train_oh,y_train_oh)
y_train_pred_oh = clf.predict(X_train_oh)
y_test_pred_oh = clf.predict(X_test_oh)
end_time = time.time()
print("Completed")
print(round((end_time-start_time)/60,2))
```

```
Completed
100.05
```

```
[22]: score_list_svc=[]
```

```
[ ]: # score_list_svc=[]
score_list.append([metrics.accuracy_score(y_train_oh, y_train_pred_oh),
                  metrics.accuracy_score(y_test_oh, y_test_pred_oh),
                  f1_score(y_train_oh, y_train_pred_oh, average="macro"),
                  f1_score(y_test_oh, y_test_pred_oh, average="macro")])
```

```
[52]: start_time = time.time()
clf = OneVsRestClassifier(SVC(kernel='linear', probability=False), n_jobs=-1)
clf.fit(X_train_tf, y_train_tf)
y_train_pred_tf = clf.predict(X_train_tf)
y_test_pred_tf = clf.predict(X_test_tf)
end_time = time.time()
print("Completed")
print(round((end_time-start_time)/60,2))
```

Completed
58.17

```
[17]: start_time = time.time()
clf = OneVsRestClassifier(SVC(kernel='linear', probability=False), n_jobs=-1)
clf.fit(X_train_we, y_train_we)
y_train_pred_we = clf.predict(X_train_we)
y_test_pred_we = clf.predict(X_test_we)
end_time = time.time()
print("Completed")
print(round(end_time-start_time,2))
```

Completed
194.55

```
[57]: score_list_svc=[]
score_list_svc=[[metrics.accuracy_score(y_train_oh, y_train_pred_oh),
                  metrics.accuracy_score(y_test_oh, y_test_pred_oh),
                  f1_score(y_train_oh, y_train_pred_oh, average="macro"),
                  f1_score(y_test_oh, y_test_pred_oh, average="macro")],
               [metrics.accuracy_score(y_train_tf, y_train_pred_tf),
                  metrics.accuracy_score(y_test_tf, y_test_pred_tf),
                  f1_score(y_train_tf, y_train_pred_tf, average="macro"),
                  f1_score(y_test_tf, y_test_pred_tf, average="macro")],
               [metrics.accuracy_score(y_train_we, y_train_pred_we),
                  metrics.accuracy_score(y_test_we, y_test_pred_we),
                  f1_score(y_train_we, y_train_pred_we, average="macro"),
                  f1_score(y_test_we, y_test_pred_we, average="macro")]]
result_svc=pd.
↳DataFrame(score_list_svc,columns=["Train_acc","Test_acc","Train_f1","Test_f1"],index=encodi
result_svc
```

[57]:		Train_acc	Test_acc	Train_f1	Test_f1
	One-hot+SVC	0.794335	0.612056	0.745194	0.403464
	TF-IDF+SVC	0.733271	0.627512	0.647086	0.383204
	Word_embedding+SVC	0.598222	0.592184	0.317670	0.300082

1.8 Feature selection

```
[26]: from sklearn import feature_selection
X_names=X_train_tf.columns
p_value_limit = 0.997
df_features = pd.DataFrame()

for cat in np.unique(y_train_tf):
    chi2, p = feature_selection.chi2(X_train_tf, y_train_tf==cat)
    df_features = df_features.append(pd.DataFrame(
        {"feature":X_names, "score":1-p, "label":cat}))
    df_features = df_features.sort_values(["label","score"],
        ascending=[True,False])
    df_features = df_features[df_features["score"]>p_value_limit]
```

```
[27]: X_train_tf_fs=X_train_tf.loc[:,df_features.feature.unique()]
X_test_tf_fs=X_test_tf.loc[:,df_features.feature.unique()]
```

```
[28]: X_train_tf_fs.shape
```

```
[28]: (18112, 206)
```

```
[29]: X_names=X_train_oh.columns
p_value_limit = 0.997
df_features = pd.DataFrame()

for cat in np.unique(y_train_oh):
    chi2, p = feature_selection.chi2(X_train_oh, y_train_oh==cat)
    df_features = df_features.append(pd.DataFrame(
        {"feature":X_names, "score":1-p, "label":cat}))
    df_features = df_features.sort_values(["label","score"],
        ascending=[True,False])
    df_features = df_features[df_features["score"]>p_value_limit]

for cat in np.unique(y_train_oh):
    print("# {}: ".format(cat))
    print(" . selected features:",len(df_features[df_features["label"]==cat]))
    print(" . top features:", ", ".
        ↪join(df_features[df_features["label"]==cat]["feature"].values[:10]))
    print(" ")
```

```
# 1:
```

```

. selected features: 285
. top features:
aw,cheap,disappoint,disintegr,excit,hideou,horribl,huge,mess,money

# 2:
. selected features: 281
. top features:
back,bad,cheap,comfort,disappoint,great,huge,model,perfect,return

# 3:
. selected features: 315
. top features:
back,bad,comfort,compliment,disappoint,excit,great,howev,jean,much

# 4:
. selected features: 137
. top features: bit,complaint,keep,littl,run,star,perfect,though,overal,nice

# 5:
. selected features: 577
. top features: absolut,amaz,area,arm,awkward,back,bad,beauti,big,boot

```

```
[31]: X_train_oh_fs=X_train_oh.loc[:,df_features.feature.unique()]
      X_test_oh_fs=X_test_oh.loc[:,df_features.feature.unique()]

```

```
[32]: X_train_oh_fs.shape

```

```
[32]: (18112, 836)
```

1.8.1 Retain on Neural Network (Keras)

```

[72]: encoding_name=["One-hot+NN+FS","TF-IDF+NN+FS"]
      data_pool=[[X_train_oh_fs.to_numpy(),X_test_oh_fs.
      ↪to_numpy(),y_train_oh,y_test_oh],
      [X_train_tf_fs.to_numpy(),X_test_tf_fs.
      ↪to_numpy(),y_train_tf,y_test_tf]]

      score_list=[]

      for data in data_pool:
          clf=NN_fit(data[0],data[2])
          # train_score=clf.evaluate(data[0],pd.get_dummies(data[2]))[1]
          # test_score=clf.evaluate(data[1],pd.get_dummies(data[3]))[1]
          # score_list.append([train_score,test_score])
          y_train_pred_nn=clf.predict_classes(data[0])+1
          y_test_pred_nn=clf.predict_classes(data[1])+1

```



```

        score_list.append([metrics.accuracy_score(data[2], y_train_pred_nn),
                           metrics.accuracy_score(data[3], y_test_pred_nn),
                           f1_score(data[2], y_train_pred_nn, average="macro"),
                           f1_score(data[3], y_test_pred_nn, average="macro")])

result_nn_fs=pd.
↳DataFrame(score_list,columns=["Train_acc","Test_acc","Train_f1","Test_f1"],index=encoding_n
result_nn_fs

```

Epoch 1/100

453/453 [=====] - 1s 1ms/step - loss: 1.2319 -
accuracy: 0.5268 - val_loss: 0.9058 - val_accuracy: 0.6243

Epoch 2/100

453/453 [=====] - 0s 1ms/step - loss: 0.8731 -
accuracy: 0.6374 - val_loss: 0.8751 - val_accuracy: 0.6392

Epoch 3/100

453/453 [=====] - 0s 1ms/step - loss: 0.8075 -
accuracy: 0.6687 - val_loss: 0.8736 - val_accuracy: 0.6420

Epoch 4/100

453/453 [=====] - 0s 987us/step - loss: 0.8011 -
accuracy: 0.6699 - val_loss: 0.8764 - val_accuracy: 0.6404

Epoch 5/100

453/453 [=====] - 0s 978us/step - loss: 0.7722 -
accuracy: 0.6856 - val_loss: 0.8771 - val_accuracy: 0.6420

Epoch 6/100

453/453 [=====] - 0s 980us/step - loss: 0.7500 -
accuracy: 0.7001 - val_loss: 0.8857 - val_accuracy: 0.6409

Epoch 1/100

453/453 [=====] - 1s 1ms/step - loss: 1.4440 -
accuracy: 0.5028 - val_loss: 1.0281 - val_accuracy: 0.5702

Epoch 2/100

453/453 [=====] - 0s 851us/step - loss: 0.9845 -
accuracy: 0.5946 - val_loss: 0.9300 - val_accuracy: 0.6133

Epoch 3/100

453/453 [=====] - 0s 857us/step - loss: 0.9307 -
accuracy: 0.6151 - val_loss: 0.9074 - val_accuracy: 0.6241

Epoch 4/100

453/453 [=====] - 0s 857us/step - loss: 0.8950 -
accuracy: 0.6323 - val_loss: 0.9004 - val_accuracy: 0.6252

Epoch 5/100

453/453 [=====] - 0s 859us/step - loss: 0.8842 -
accuracy: 0.6345 - val_loss: 0.8962 - val_accuracy: 0.6257

Epoch 6/100

453/453 [=====] - 0s 858us/step - loss: 0.8756 -
accuracy: 0.6385 - val_loss: 0.8949 - val_accuracy: 0.6279

Epoch 7/100

453/453 [=====] - 0s 861us/step - loss: 0.8599 -

```

accuracy: 0.6427 - val_loss: 0.8937 - val_accuracy: 0.6288
Epoch 8/100
453/453 [=====] - 0s 968us/step - loss: 0.8763 -
accuracy: 0.6335 - val_loss: 0.8946 - val_accuracy: 0.6288
Epoch 9/100
453/453 [=====] - 0s 935us/step - loss: 0.8649 -
accuracy: 0.6421 - val_loss: 0.8913 - val_accuracy: 0.6293
Epoch 10/100
453/453 [=====] - 0s 936us/step - loss: 0.8547 -
accuracy: 0.6508 - val_loss: 0.8951 - val_accuracy: 0.6271
Epoch 11/100
453/453 [=====] - 0s 910us/step - loss: 0.8691 -
accuracy: 0.6394 - val_loss: 0.8933 - val_accuracy: 0.6268
Epoch 12/100
453/453 [=====] - 0s 992us/step - loss: 0.8557 -
accuracy: 0.6531 - val_loss: 0.8968 - val_accuracy: 0.6227

```

[72]:		Train_acc	Test_acc	Train_f1	Test_f1
	One-hot+NN+FS	0.698929	0.631486	0.545134	0.417691
	TF-IDF+NN+FS	0.643883	0.619121	0.458443	0.402997

1.8.2 Retrain on linear SVC

```
[41]: score_list_svc_fs=[]
```

```

[42]: start_time = time.time()
      clf = OneVsRestClassifier(SVC(kernel='linear', probability=False), n_jobs=-1)
      clf.fit(X_train_oh_fs, y_train_oh)
      y_train_pred_oh_fs = clf.predict(X_train_oh_fs)
      y_test_pred_oh_fs = clf.predict(X_test_oh_fs)
      end_time = time.time()
      print("Completed")
      print(round(end_time-start_time,2))

```

```

Completed
1753.59

```

```

[43]: score_list_svc_fs.append([metrics.accuracy_score(y_train_oh,
↪y_train_pred_oh_fs),
                                metrics.accuracy_score(y_test_oh, y_test_pred_oh_fs),
                                f1_score(y_train_oh, y_train_pred_oh_fs, average="macro"),
                                f1_score(y_test_oh, y_test_pred_oh_fs, average="macro")])

```

```

[44]: clf = OneVsRestClassifier(SVC(kernel='linear', probability=False), n_jobs=-1)
      clf.fit(X_train_tf_fs, y_train_tf)
      y_train_pred_tf_fs = clf.predict(X_train_tf_fs)
      y_test_pred_tf_fs = clf.predict(X_test_tf_fs)

```

```
[45]: # score_list_svc=[]
score_list_svc_fs.append([metrics.accuracy_score(y_train_tf,
↪y_train_pred_tf_fs),
                        metrics.accuracy_score(y_test_tf, y_test_pred_tf_fs),
                        f1_score(y_train_tf, y_train_pred_tf_fs, average="macro"),
                        f1_score(y_test_tf, y_test_pred_tf_fs, average="macro")])
```

```
[75]: result_svc_fs=pd.
↪DataFrame(score_list_svc_fs,columns=["Train_acc","Test_acc","Train_f1","Test_f1"],index=["0",
result_svc_fs
```

```
[75]:
```

	Train_acc	Test_acc	Train_f1	Test_f1
One-hot+SVC+FS	0.676734	0.609848	0.533671	0.366183
TF-IDF+SVC+FS	0.615007	0.603444	0.380943	0.321392

```
[81]: result_nn.append(result_svc).round(4)
```

```
[81]:
```

	Train_acc	Test_acc	Train_f1	Test_f1
One-hot+NN	0.7554	0.6182	0.6366	0.4437
TF-IDF+NN	0.7419	0.6125	0.6210	0.4212
Word_embedding+NN	0.6136	0.6083	0.3560	0.3469
One-hot+SVC	0.7943	0.6121	0.7452	0.4035
TF-IDF+SVC	0.7333	0.6275	0.6471	0.3832
Word_embedding+SVC	0.5982	0.5922	0.3177	0.3001

```
[84]: result_nn_fs.append(result_svc_fs).round(4)
```

```
[84]:
```

	Train_acc	Test_acc	Train_f1	Test_f1
One-hot+NN+FS	0.6989	0.6315	0.5451	0.4177
TF-IDF+NN+FS	0.6439	0.6191	0.4584	0.4030
One-hot+SVC+FS	0.6767	0.6098	0.5337	0.3662
TF-IDF+SVC+FS	0.6150	0.6034	0.3809	0.3214

```
[83]: result_nn.append(result_svc).append(result_svc_fs).append(result_nn_fs).round(4)
```

```
[83]:
```

	Train_acc	Test_acc	Train_f1	Test_f1
One-hot+NN	0.7554	0.6182	0.6366	0.4437
TF-IDF+NN	0.7419	0.6125	0.6210	0.4212
Word_embedding+NN	0.6136	0.6083	0.3560	0.3469
One-hot+SVC	0.7943	0.6121	0.7452	0.4035
TF-IDF+SVC	0.7333	0.6275	0.6471	0.3832
Word_embedding+SVC	0.5982	0.5922	0.3177	0.3001
One-hot+SVC+FS	0.6767	0.6098	0.5337	0.3662
TF-IDF+SVC+FS	0.6150	0.6034	0.3809	0.3214
One-hot+NN+FS	0.6989	0.6315	0.5451	0.4177
TF-IDF+NN+FS	0.6439	0.6191	0.4584	0.4030

1.9 Recommender System

```
[82]: reviews=pd.read_csv("data/Womens Clothing E-Commerce Reviews.csv")
      df = reviews
```

```
[83]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23486 entries, 0 to 23485
Data columns (total 11 columns):
Unnamed: 0                23486 non-null int64
Clothing ID               23486 non-null int64
Age                       23486 non-null int64
Title                    19676 non-null object
Review Text              22641 non-null object
Rating                   23486 non-null int64
Recommended IND          23486 non-null int64
Positive Feedback Count  23486 non-null int64
Division Name            23472 non-null object
Department Name          23472 non-null object
Class Name               23472 non-null object
dtypes: int64(6), object(5)
memory usage: 2.0+ MB
```

```
[84]: df[df['Division Name'].isnull()]
```

```
[84]:
```

	Unnamed: 0	Clothing ID	Age	Title \
9444	9444	72	25	My favorite socks!!!
13767	13767	492	23	So soft!
13768	13768	492	49	Wardrobe staple
13787	13787	492	48	NaN
16216	16216	152	36	Warm and cozy
16221	16221	152	37	Love!
16223	16223	152	39	"long and warm"
18626	18626	184	34	Nubby footless tights
18671	18671	184	54	New workhorse
20088	20088	772	50	Comfy sweatshirt!
21532	21532	665	43	So worth it!
22997	22997	136	47	Charcoal, pale gray, a bit of silver!
23006	23006	136	33	Cute itsy socks
23011	23011	136	36	Super socks

	Review Text	Rating \
9444	I never write reviews, but these socks are so ...	5
13767	I just love this hoodie! it is so soft and com...	5
13768	Love this hoodie. so soft and goes with everyt...	5
13787	NaN	5

16216	Just what i was looking for. soft, cozy and warm.	5
16221	I am loving these. they are quite long but are...	5
16223	These leg warmers are perfect for me. they are...	5
18626	These are amazing quality. i agree, size up to...	5
18671	These tights are amazing! if i care for them w...	5
20088	This sweatshirt is really nice! it's oversize...	5
21532	Got these on sale...absolutely love eberjey! f...	5
22997	These socks are soft and comfortable, and they...	5
23006	Love polkadots, love sparkle. these little soc...	5
23011	I love these little socks ? and the dots spark...	5

	Recommended IND	Positive Feedback Count	Division Name	Department Name	\
9444	1	0	NaN	NaN	
13767	1	1	NaN	NaN	
13768	1	0	NaN	NaN	
13787	1	0	NaN	NaN	
16216	1	0	NaN	NaN	
16221	1	0	NaN	NaN	
16223	1	0	NaN	NaN	
18626	1	5	NaN	NaN	
18671	1	0	NaN	NaN	
20088	1	0	NaN	NaN	
21532	1	0	NaN	NaN	
22997	1	1	NaN	NaN	
23006	1	0	NaN	NaN	
23011	1	0	NaN	NaN	

	Class Name
9444	NaN
13767	NaN
13768	NaN
13787	NaN
16216	NaN
16221	NaN
16223	NaN
18626	NaN
18671	NaN
20088	NaN
21532	NaN
22997	NaN
23006	NaN
23011	NaN

```
[85]: df_filter = df.dropna(subset = ['Division Name', 'Department Name', 'Class_
↳Name'])
```

```
[86]: df_filter.head()
```

```
[86]: Unnamed: 0 Clothing ID Age Title \
0 0 767 33 NaN
1 1 1080 34 NaN
2 2 1077 60 Some major design flaws
3 3 1049 50 My favorite buy!
4 4 847 47 Flattering shirt

Review Text Rating Recommended IND \
0 Absolutely wonderful - silky and sexy and comf... 4 1
1 Love this dress! it's sooo pretty. i happene... 5 1
2 I had such high hopes for this dress and reall... 3 0
3 I love, love, love this jumpsuit. it's fun, fl... 5 1
4 This shirt is very flattering to all due to th... 5 1

Positive Feedback Count Division Name Department Name Class Name
0 0 Initmates Intimate Intimates
1 4 General Dresses Dresses
2 0 General Dresses Dresses
3 0 General Petite Bottoms Pants
4 6 General Tops Blouses
```

```
[87]: bins = [18, 30, 40, 50, 60, 70, 120]
labels = ['18-29', '30-39', '40-49', '50-59', '60-69', '70+']
df_filter['Age_group'] = pd.cut(df_filter.Age, bins, labels=labels,
    include_lowest = True, right = False)
```

```
[88]: df_filter['Group_rating'] = df_filter.groupby(['Clothing_ID', 'Age_group'])[['Rating']].transform('mean')
```

```
[89]: df_filter['Group_recommend_prop'] = df_filter.groupby(['Clothing_ID', 'Age_group'])[['Recommended IND']].transform('mean')
```

```
[90]: Age_groups = pd.get_dummies(df_filter.Age_group)
Age_groups[:10]
```

```
[90]: Age_group 18-29 30-39 40-49 50-59 60-69 70+
0 0 1 0 0 0 0
1 0 1 0 0 0 0
2 0 0 0 0 1 0
3 0 0 0 1 0 0
4 0 0 1 0 0 0
5 0 0 1 0 0 0
6 0 1 0 0 0 0
7 0 1 0 0 0 0
8 1 0 0 0 0 0
9 0 1 0 0 0 0
```

```

[91]: df_filter = pd.concat([df_filter, Age_groups], axis = 1)

[92]: df_filter['Rating(18-29)'] = df_filter['Rating'] * df_filter['18-29']
df_filter['Rating(30-39)'] = df_filter['Rating'] * df_filter['30-39']
df_filter['Rating(40-49)'] = df_filter['Rating'] * df_filter['40-49']
df_filter['Rating(50-59)'] = df_filter['Rating'] * df_filter['50-59']
df_filter['Rating(60-69)'] = df_filter['Rating'] * df_filter['60-69']
df_filter['Rating(70+)'] = df_filter['Rating'] * df_filter['70+']
df_filter['Recommend_prop(18-29)'] = df_filter['Recommended IND'] *
    ↪df_filter['18-29']
df_filter['Recommend_prop(30-39)'] = df_filter['Recommended IND'] *
    ↪df_filter['30-39']
df_filter['Recommend_prop(40-49)'] = df_filter['Recommended IND'] *
    ↪df_filter['40-49']
df_filter['Recommend_prop(50-59)'] = df_filter['Recommended IND'] *
    ↪df_filter['50-59']
df_filter['Recommend_prop(60-69)'] = df_filter['Recommended IND'] *
    ↪df_filter['60-69']
df_filter['Recommend_prop(70+)'] = df_filter['Recommended IND'] *
    ↪df_filter['70+']

[93]: product_vec = df_filter.groupby('Clothing ID')[['18-29', '30-39', '40-49',
    '50-59', '60-69', '70+', 'Rating(18-29)', 'Rating(30-39)',
    'Rating(40-49)', 'Rating(50-59)', 'Rating(60-69)', 'Rating(70+)',
    'Recommend_prop(18-29)', 'Recommend_prop(30-39)',
    'Recommend_prop(40-49)', 'Recommend_prop(50-59)',
    'Recommend_prop(60-69)', 'Recommend_prop(70+)']].agg('sum').reset_index()

[94]: for i in ['18-29', '30-39', '40-49', '50-59', '60-69', '70+']:
    col_rating = 'Rating('+i+')'
    col_recom = 'Recommend_prop('+i+')'
    product_vec[col_rating] = round(product_vec[col_rating] / product_vec[i],3)
    product_vec[col_recom] = round(product_vec[col_recom] / product_vec[i],3)

[95]: product_vec = product_vec.fillna(0)

[96]: division = pd.pivot_table(df_filter[["Clothing ID", "Division_
    ↪Name"]], index="Clothing ID", columns="Division Name", values="Clothing_
    ↪ID", aggfunc=lambda x:1 if len(x)>0 else 0).fillna(0).reset_index()

[97]: division.index=product_vec.index

[98]: product_vec = product_vec.merge(division, on='Clothing ID', how='left')

[99]: # scale the number of each age group

```

```

product_vec['total'] =
↳product_vec['18-29']+product_vec['30-39']+product_vec['40-49']+product_vec['50-59']+product
product_vec['18-29'] = round(product_vec['18-29']/product_vec['total'],3)
product_vec['30-39'] = round(product_vec['30-39']/product_vec['total'],3)
product_vec['40-49'] = round(product_vec['40-49']/product_vec['total'],3)
product_vec['50-59'] = round(product_vec['50-59']/product_vec['total'],3)
product_vec['60-69'] = round(product_vec['60-69']/product_vec['total'],3)
product_vec['70+'] = round(product_vec['70+']/product_vec['total'],3)
del product_vec['total']
product_vec

```

```

[99]:      Clothing ID  18-29  30-39  40-49  50-59  60-69  70+  Rating(18-29) \
0          0      1.000  0.000  0.000  0.000  0.000  0.0      5.0
1          1      0.333  0.333  0.000  0.333  0.000  0.0      2.0
2          2      1.000  0.000  0.000  0.000  0.000  0.0      4.0
3          3      0.000  1.000  0.000  0.000  0.000  0.0      0.0
4          4      1.000  0.000  0.000  0.000  0.000  0.0      5.0
...
1194      1201  0.000  1.000  0.000  0.000  0.000  0.0      0.0
1195      1202  0.111  0.333  0.333  0.111  0.111  0.0      4.0
1196      1203  0.182  0.455  0.182  0.182  0.000  0.0      4.0
1197      1204  0.000  1.000  0.000  0.000  0.000  0.0      0.0
1198      1205  0.000  0.500  0.000  0.000  0.500  0.0      0.0

```

```

      Rating(30-39)  Rating(40-49)  ...  Rating(70+)  Recommend_prop(18-29) \
0          0.0          0.000  ...          0.0          1.0
1          5.0          0.000  ...          0.0          0.0
2          0.0          0.000  ...          0.0          1.0
3          5.0          0.000  ...          0.0          0.0
4          0.0          0.000  ...          0.0          1.0
...
1194          4.0          0.000  ...          0.0          0.0
1195          5.0          3.333  ...          0.0          1.0
1196          4.8          5.000  ...          0.0          0.5
1197          4.5          0.000  ...          0.0          0.0
1198          5.0          0.000  ...          0.0          0.0

```

```

      Recommend_prop(30-39)  Recommend_prop(40-49)  Recommend_prop(50-59) \
0          0.0          0.0          0.0
1          1.0          0.0          1.0
2          0.0          0.0          0.0
3          1.0          0.0          0.0
4          0.0          0.0          0.0
...
1194          1.0          0.0          0.0
1195          1.0          1.0          1.0
1196          1.0          1.0          0.5

```


1197	1.0	0.0	0.0
1198	1.0	0.0	0.0

	Recommend_prop(60-69)	Recommend_prop(70+)	General	General Petite \
0	0.0	0.0	1.0	0.0
1	0.0	0.0	0.0	0.0
2	0.0	0.0	1.0	0.0
3	0.0	0.0	1.0	0.0
4	0.0	0.0	1.0	0.0
...
1194	0.0	0.0	1.0	0.0
1195	1.0	0.0	0.0	1.0
1196	0.0	0.0	0.0	0.0
1197	0.0	0.0	0.0	1.0
1198	1.0	0.0	1.0	0.0

	Initmates
0	0.0
1	1.0
2	0.0
3	0.0
4	0.0
...	...
1194	0.0
1195	0.0
1196	1.0
1197	0.0
1198	0.0

[1199 rows x 22 columns]

Calculate similarity

```
[100]: product_vec = product_vec.set_index('Clothing ID')
```

```
[101]: general_prod = product_vec.loc[product_vec.General == 1, "18-29":
↳ "Recommend_prop(70+)"]
```

```
[102]: general_petite_prod = product_vec.loc[product_vec['General Petite'] == 1,
↳ "18-29": "Recommend_prop(70+)"]
```

```
[103]: initmates_prod = product_vec.loc[product_vec.Initmates == 1, "18-29":
↳ "Recommend_prop(70+)"]
```

```
[104]: from sklearn.neighbors import NearestNeighbors
neigh = NearestNeighbors(metric='cosine', algorithm='brute', n_neighbors=6)
```

```
[105]: neigh.fit(general_prod)
neighbors_general = neigh.kneighbors(general_prod, return_distance = False)
```

```
[106]: # fit models for general product
neigh.fit(general_prod)
neighbors_general = neigh.kneighbors(general_prod, return_distance = False)
neighbors_general = pd.DataFrame(neighbors_general, index = general_prod.index)
neighbors_general.columns = ['Top1', 'Top2', 'Top3', 'Top4', 'Top5', 'Top6']
neighbors_general['Top1'] = neighbors_general.index[neighbors_general.Top1]
neighbors_general['Top2'] = neighbors_general.index[neighbors_general.Top2]
neighbors_general['Top3'] = neighbors_general.index[neighbors_general.Top3]
neighbors_general['Top4'] = neighbors_general.index[neighbors_general.Top4]
neighbors_general['Top5'] = neighbors_general.index[neighbors_general.Top5]
neighbors_general['Top6'] = neighbors_general.index[neighbors_general.Top6]
neighbors_general
```

```
[106]:
```

	Top1	Top2	Top3	Top4	Top5	Top6
Clothing ID						
0	54	958	36	12	75	575
2	19	989	633	45	526	2
3	541	485	25	548	1043	299
4	54	958	36	12	75	575
5	541	485	25	548	1043	299
...
1197	1197	551	177	882	1118	1137
1198	1198	550	588	1039	1113	925
1200	438	429	1200	53	1127	982
1201	20	587	9	1186	26	13
1205	1205	543	979	1120	1145	542

[454 rows x 6 columns]

```
[107]: # fit models for general petite product
neigh.fit(general_petite_prod)
neighbors_petite = neigh.kneighbors(general_petite_prod, return_distance = 
    ↪ False)
neighbors_petite = pd.DataFrame(neighbors_petite, index = general_petite_prod.
    ↪ index)
neighbors_petite.columns = ['Top1', 'Top2', 'Top3', 'Top4', 'Top5', 'Top6']
neighbors_petite['Top1'] = neighbors_petite.index[neighbors_petite.Top1]
neighbors_petite['Top2'] = neighbors_petite.index[neighbors_petite.Top2]
neighbors_petite['Top3'] = neighbors_petite.index[neighbors_petite.Top3]
neighbors_petite['Top4'] = neighbors_petite.index[neighbors_petite.Top4]
neighbors_petite['Top5'] = neighbors_petite.index[neighbors_petite.Top5]
neighbors_petite['Top6'] = neighbors_petite.index[neighbors_petite.Top6]
neighbors_petite
```

```
[107]:
```

	Top1	Top2	Top3	Top4	Top5	Top6
Clothing ID						
18	18	421	1179	615	703	23
21	340	1014	1124	21	756	1007
23	639	23	703	957	615	1179
29	29	1196	356	849	561	1140
30	30	1199	771	473	1204	1149
...
1190	1190	1067	490	649	959	1184
1196	29	1196	356	849	561	1140
1199	30	1199	771	473	1204	1149
1202	1202	1042	1131	876	1051	839
1204	473	1204	1149	589	1029	1182

[308 rows x 6 columns]

```
[108]: # fit models for initmates product
neigh.fit(initmates_prod)
neighbors_initmates = neigh.kneighbors(initmates_prod, return_distance = False)
neighbors_initmates = pd.DataFrame(neighbors_initmates, index = initmates_prod.
    ↪ index)
neighbors_initmates.columns = ['Top1', 'Top2', 'Top3', 'Top4', 'Top5', 'Top6']
neighbors_initmates['Top1'] = neighbors_initmates.index[neighbors_initmates.
    ↪ Top1]
neighbors_initmates['Top2'] = neighbors_initmates.index[neighbors_initmates.
    ↪ Top2]
neighbors_initmates['Top3'] = neighbors_initmates.index[neighbors_initmates.
    ↪ Top3]
neighbors_initmates['Top4'] = neighbors_initmates.index[neighbors_initmates.
    ↪ Top4]
neighbors_initmates['Top5'] = neighbors_initmates.index[neighbors_initmates.
    ↪ Top5]
neighbors_initmates['Top6'] = neighbors_initmates.index[neighbors_initmates.
    ↪ Top6]
neighbors_initmates
```

```
[108]:
```

	Top1	Top2	Top3	Top4	Top5	Top6
Clothing ID						
1	1	684	441	163	268	231
8	8	225	360	242	107	442
10	140	126	637	655	149	544
14	98	707	519	729	517	716
15	15	625	706	1167	741	259
...
1175	114	87	218	410	207	434
1176	219	432	1176	305	96	210
1177	140	126	637	655	149	544

```
[633 rows x 6 columns]
```

```
[110]: neighbors = pd.DataFrame(neighbors)
neighbors.columns = ['Top1', 'Top2', 'Top3', 'Top4', 'Top5', 'Top6']
neighbors
```

```
[1199 rows x 6 columns]
```

```
[112]: product_vec[product_vec.index.isin([18,421,1179,615,703,23])]
```

60

Clothing ID				...	
18	0.0	1.0	0.0	...	0.0
23	0.0	5.0	0.0	...	0.0
421	0.0	1.0	0.0	...	0.0
615	0.0	4.0	0.0	...	0.0
703	0.0	5.0	0.0	...	0.0
1179	0.0	4.0	0.0	...	0.0

	Recommend_prop(18-29)	Recommend_prop(30-39)	\
Clothing ID			
18	0.0	0.0	
23	0.0	0.0	
421	0.0	0.0	
615	0.0	0.0	
703	0.0	0.0	
1179	0.0	0.0	

	Recommend_prop(40-49)	Recommend_prop(50-59)	\
Clothing ID			
18	0.0	0.0	
23	1.0	0.0	
421	0.0	0.0	
615	1.0	0.0	
703	1.0	0.0	
1179	1.0	0.0	

	Recommend_prop(60-69)	Recommend_prop(70+)	General	\
Clothing ID				
18	0.0	0.0	0.0	
23	0.0	0.0	0.0	
421	0.0	0.0	0.0	
615	0.0	0.0	0.0	
703	0.0	0.0	0.0	
1179	0.0	0.0	0.0	

	General Petite	Initmates
Clothing ID		
18	1.0	0.0
23	1.0	0.0
421	1.0	0.0
615	1.0	0.0
703	1.0	0.0
1179	1.0	0.0

[6 rows x 21 columns]

1.10 Testing Recommendation System on Another Dataset

```
[113]: df_book = pd.read_csv("data/Books.csv")
df_rating = pd.read_csv("data/Ratings.csv")
df_user = pd.read_csv("data/Users.csv")
```

```
[114]: df_book = df_book.iloc[:100000,:]
df_rating = df_rating.iloc[:100000,:]
df_user = df_user.iloc[:100000,:]
df_user = df_user.drop(['Location'],axis = 1)
df_rating = df_rating.iloc[:, 0:3]
df_user
```

```
[114]:
```

	User-ID	Age
0	1	NaN
1	2	18.0
2	3	NaN
3	4	17.0
4	5	NaN
...
99995	99996	43.0
99996	99997	33.0
99997	99998	22.0
99998	99999	46.0
99999	100000	43.0

[100000 rows x 2 columns]

```
[115]: df_book_rating = pd.merge(df_book, df_rating, on="ISBN")
df_all = pd.merge(df_book_rating, df_user, on="User-ID")
```

```
[116]: df_all = df_all.dropna(subset = ['Age'])
df_all = df_all.drop(['Image-URL-S', 'Image-URL-M', 'Image-URL-L'],axis = 1)
df_all.rename(columns = {'Book-Rating':'Rating'}, inplace = True)
df_all['Year-Of-Publication'] = df_all['Year-Of-Publication'].astype('category')
#i = df_all[((df_all['Year-Of-Publication'] == 0) | (  
→df_all['Year-Of-Publication'] == 2030))].index
#df_all = df_all.drop(i,axis = 0)
df_all
```

```
[116]:
```

	ISBN	Book-Title \
0	0195153448	Classical Mythology
18	0002005018	Clara Callan
19	0786868716	The Five People You Meet in Heaven
20	0151008116	Life of Pi
21	0671021001	She's Come Undone (Oprah's Book Club)
...

74543	0140327592		Matilda
74545	1400049520	Slander : Liberal Lies About the American Right	
74546	8440655193	Avicena, O La Ruta de Isfahan	
74547	0886771404	The Shapechangers	
74548	0312422288	Sellelevision: A Novel	

	Book-Author	Year-Of-Publication	Publisher \
0	Mark P. O. Morford	2002	Oxford University Press
18	Richard Bruce Wright	2001	HarperFlamingo Canada
19	Mitch Albom	2003	Hyperion
20	Yann Martel	2002	Harcourt
21	Wally Lamb	1998	Pocket
...
74543	Roald Dahl	1990	Viking Penguin Inc
74545	ANN COULTER	2003	Three Rivers Press
74546	Gilbert Sinove	1997	Ediciones B
74547	Jennifer Roberson	1992	Daw Books
74548	Augusten Burroughs	2003	Picador USA

	User-ID	Rating	Age
0	2	0	18.0
18	11400	0	49.0
19	11400	9	49.0
20	11400	6	49.0
21	11400	0	49.0
...
74543	16923	6	16.0
74545	4806	9	29.0
74546	9131	0	37.0
74547	5775	6	34.0
74548	10594	4	57.0

[49156 rows x 8 columns]

```
[117]: bins = [1920,1930, 1940, 1950, 1960, 1970, 1980,1990,2000,2010,2020]
labels = ['1920-1930','1930-1940', '1940-1950', '1950-1960','1960-1970',
↪ '1970-1980', '1980-1990', '1990-2000','2000-2010','2010-2020']
df_all['Year_group'] = pd.cut(df_all['Year-Of-Publication'], bins,
↪ labels=labels, include_lowest = True, right = False)
df_all = df_all.dropna(subset = ['Year_group'])

bins = [1,18, 30, 40, 50, 60, 70, 120]
labels = ['1-17','18-29', '30-39', '40-49', '50-59', '60-69', '70+']
df_all['Age_group'] = pd.cut(df_all.Age, bins, labels=labels, include_lowest =
↪ True, right = False)
```

```
df_all['Group_rating'] = df_all.groupby(['ISBN', 'Age_group'])[['Rating']].
    ↪transform('mean')
df_all['Group_Year_rating'] = df_all.groupby(['Year_group'])[['Rating']].
    ↪transform('mean')
df_all
```

```
[117]:
```

	ISBN	Book-Title \
0	0195153448	Classical Mythology
18	0002005018	Clara Callan
19	0786868716	The Five People You Meet in Heaven
20	0151008116	Life of Pi
21	0671021001	She's Come Undone (Oprah's Book Club)
...
74543	0140327592	Matilda
74545	1400049520	Slander : Liberal Lies About the American Right
74546	8440655193	Avicena, O La Ruta de Isfahan
74547	0886771404	The Shapechangers
74548	0312422288	Sellelevision: A Novel

	Book-Author	Year-Of-Publication	Publisher \
0	Mark P. O. Morford	2002	Oxford University Press
18	Richard Bruce Wright	2001	HarperFlamingo Canada
19	Mitch Albom	2003	Hyperion
20	Yann Martel	2002	Harcourt
21	Wally Lamb	1998	Pocket
...
74543	Roald Dahl	1990	Viking Penguin Inc
74545	ANN COULTER	2003	Three Rivers Press
74546	Gilbert Sinove	1997	Ediciones B
74547	Jennifer Roberson	1992	Daw Books
74548	Augusten Burroughs	2003	Picador USA

	User-ID	Rating	Age	Year_group	Age_group	Group_rating \
0	2	0	18.0	2000-2010	18-29	NaN
18	11400	0	49.0	2000-2010	40-49	NaN
19	11400	9	49.0	2000-2010	40-49	NaN
20	11400	6	49.0	2000-2010	40-49	NaN
21	11400	0	49.0	1990-2000	40-49	0.0
...
74543	16923	6	16.0	1990-2000	1-17	NaN
74545	4806	9	29.0	2000-2010	18-29	NaN
74546	9131	0	37.0	1990-2000	30-39	NaN
74547	5775	6	34.0	1990-2000	30-39	NaN
74548	10594	4	57.0	2000-2010	50-59	10.0

	Group_Year_rating
0	3.222195


```

18          3.222195
19          3.222195
20          3.222195
21          3.015357
...
74543       3.015357
74545       3.222195
74546       3.015357
74547       3.015357
74548       3.222195

```

```
[48459 rows x 12 columns]
```

```
[118]: Age_groups = pd.get_dummies(df_all.Age_group)
       Year_groups = pd.get_dummies(df_all['Year_group'])
```

```
[119]: df_all = pd.concat([df_all, Age_groups, Year_groups], axis = 1)
```

```
[120]: df_all['Rating(1-17)'] = df_all['Rating'] * df_all['1-17']
df_all['Rating(18-29)'] = df_all['Rating'] * df_all['18-29']
df_all['Rating(30-39)'] = df_all['Rating'] * df_all['30-39']
df_all['Rating(40-49)'] = df_all['Rating'] * df_all['40-49']
df_all['Rating(50-59)'] = df_all['Rating'] * df_all['50-59']
df_all['Rating(60-69)'] = df_all['Rating'] * df_all['60-69']
df_all['Rating(70+)'] = df_all['Rating'] * df_all['70+']

df_all['Rating(1920-1930)'] = df_all['Rating'] * df_all['1920-1930']
df_all['Rating(1930-1940)'] = df_all['Rating'] * df_all['1930-1940']
df_all['Rating(1940-1950)'] = df_all['Rating'] * df_all['1940-1950']
df_all['Rating(1950-1960)'] = df_all['Rating'] * df_all['1950-1960']
df_all['Rating(1960-1970)'] = df_all['Rating'] * df_all['1960-1970']
df_all['Rating(1970-1980)'] = df_all['Rating'] * df_all['1970-1980']
df_all['Rating(1980-1990)'] = df_all['Rating'] * df_all['1980-1990']
df_all['Rating(1990-2000)'] = df_all['Rating'] * df_all['1990-2000']
df_all['Rating(2000-2010)'] = df_all['Rating'] * df_all['2000-2010']
df_all['Rating(2010-2020)'] = df_all['Rating'] * df_all['2010-2020']
```

```
[121]: product_vec = df_all.groupby('ISBN')[['1-17', '18-29', '30-39', '40-49',
      '50-59', '60-69', '70+', 'Rating(1-17)', 'Rating(18-29)', 'Rating(30-39)',
      'Rating(40-49)', 'Rating(50-59)', 'Rating(60-69)', 'Rating(70+)',
      ↪ '1920-1930',
      '1930-1940', '1940-1950', '1950-1960', '1960-1970', '1970-1980',
      '1980-1990', '1990-2000', '2000-2010', '2010-2020',
      ↪ 'Rating(1920-1930)', 'Rating(1930-1940)',
      ↪
      ↪ 'Rating(1940-1950)', 'Rating(1950-1960)', 'Rating(1960-1970)', 'Rating(1970-1980)', 'Rating(1980-1990)',
      'Rating(1990-2000)', 'Rating(2000-2010)', 'Rating(2010-2020)']
```

```
]].agg('sum').reset_index()
```

```
[122]: for i in ['1-17', '18-29', '30-39', '40-49', '50-59', '60-69', '70+']:
        col_rating = 'Rating('+i+')'
        product_vec[col_rating] = round(product_vec[col_rating] / product_vec[i],3)

for i in ['1920-1930', '1930-1940', '1940-1950', '1950-1960', '1960-1970',
        ↪ '1970-1980',
        '1980-1990', '1990-2000', '2000-2010', '2010-2020']:
        col_rating = 'Rating('+i+')'
        product_vec[col_rating] = round(product_vec[col_rating] / product_vec[i],3)
```

```
[123]: product_vec = product_vec.fillna(0)
        product_vec
```

```
[123]:
```

	ISBN	1-17	18-29	30-39	40-49	50-59	60-69	70+	Rating(1-17)	\
0	0001047973	0	1	0	0	0	0	0	0.0	
1	0001360469	0	0	0	0	0	1	0	0.0	
2	0001372564	0	0	0	1	0	0	0	0.0	
3	0001374869	0	0	0	0	0	1	0	0.0	
4	0001939203	0	0	1	0	0	0	0	0.0	
...
32450	9997405307	0	1	0	0	0	0	0	0.0	
32451	9997522052	0	0	1	0	0	0	0	0.0	
32452	9999364497	0	0	1	0	0	0	0	0.0	
32453	B00007MF56	0	1	0	0	0	0	0	0.0	
32454	B0000DAPP1	0	0	1	0	0	0	0	0.0	

	Rating(18-29)	...	Rating(1920-1930)	Rating(1930-1940)	\
0	9.0	...	0.0	0.0	
1	0.0	...	0.0	0.0	
2	0.0	...	0.0	0.0	
3	0.0	...	0.0	0.0	
4	0.0	...	0.0	0.0	
...
32450	0.0	...	0.0	0.0	
32451	0.0	...	0.0	0.0	
32452	0.0	...	0.0	0.0	
32453	9.0	...	0.0	0.0	
32454	0.0	...	0.0	0.0	

	Rating(1940-1950)	Rating(1950-1960)	Rating(1960-1970)	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

```

...
32450      0.0      0.0      0.0
32451      0.0      0.0      0.0
32452      0.0      0.0      0.0
32453      0.0      0.0      0.0
32454      0.0      0.0      0.0

      Rating(1970-1980)  Rating(1980-1990)  Rating(1990-2000)  \
0      0.0      0.0      9.0
1      0.0      0.0      10.0
2      0.0      0.0      0.0
3      0.0      0.0      10.0
4      0.0      0.0      8.0
...
32450      0.0      0.0      0.0
32451      0.0      0.0      0.0
32452      0.0      0.0      0.0
32453      0.0      0.0      0.0
32454      0.0      0.0      0.0

      Rating(2000-2010)  Rating(2010-2020)
0      0.0      0.0
1      0.0      0.0
2      0.0      0.0
3      0.0      0.0
4      0.0      0.0
...
32450      0.0      0.0
32451      0.0      0.0
32452      0.0      0.0
32453      9.0      0.0
32454      5.0      0.0

```

[32455 rows x 35 columns]

```

[124]: product_vec['total_age'] = product_vec['1-17'] +
      ↪product_vec['18-29']+product_vec['30-39']+product_vec['40-49']+product_vec['50-59']+product
product_vec['1-17'] = round(product_vec['1-17']/product_vec['total_age'],3)
product_vec['18-29'] = round(product_vec['18-29']/product_vec['total_age'],3)
product_vec['30-39'] = round(product_vec['30-39']/product_vec['total_age'],3)
product_vec['40-49'] = round(product_vec['40-49']/product_vec['total_age'],3)
product_vec['50-59'] = round(product_vec['50-59']/product_vec['total_age'],3)
product_vec['60-69'] = round(product_vec['60-69']/product_vec['total_age'],3)
product_vec['70+'] = round(product_vec['70+']/product_vec['total_age'],3)
del product_vec['total_age']

```

```

product_vec['total_year'] = product_vec['1920-1930'] +
    ↳product_vec['1930-1940']+product_vec['1940-1950']+product_vec['1950-1960']+product_vec['1960-1970']
    ↳+ product_vec['1990-2000'] + product_vec['2000-2010'] +
    ↳product_vec['2010-2020']

product_vec['1920-1930'] = round(product_vec['1920-1930']/
    ↳product_vec['total_year'],3)
product_vec['1930-1940'] = round(product_vec['1930-1940']/
    ↳product_vec['total_year'],3)
product_vec['1940-1950'] = round(product_vec['1940-1950']/
    ↳product_vec['total_year'],3)
product_vec['1950-1960'] = round(product_vec['1950-1960']/
    ↳product_vec['total_year'],3)
product_vec['1960-1970'] = round(product_vec['1960-1970']/
    ↳product_vec['total_year'],3)
product_vec['1970-1980'] = round(product_vec['1970-1980']/
    ↳product_vec['total_year'],3)
product_vec['1980-1990'] = round(product_vec['1980-1990']/
    ↳product_vec['total_year'],3)
product_vec['1990-2000'] = round(product_vec['1990-2000']/
    ↳product_vec['total_year'],3)
product_vec['2000-2010'] = round(product_vec['2000-2010']/
    ↳product_vec['total_year'],3)
product_vec['2010-2020'] = round(product_vec['2010-2020']/
    ↳product_vec['total_year'],3)
del product_vec['total_year']
product_vec

```

```

[124]:
      ISBN 1-17 18-29 30-39 40-49 50-59 60-69 70+ Rating(1-17) \
0      0001047973 0.0 1.0 0.0 0.0 0.0 0.0 0.0 0.0
1      0001360469 0.0 0.0 0.0 0.0 0.0 1.0 0.0 0.0
2      0001372564 0.0 0.0 0.0 1.0 0.0 0.0 0.0 0.0
3      0001374869 0.0 0.0 0.0 0.0 0.0 1.0 0.0 0.0
4      0001939203 0.0 0.0 1.0 0.0 0.0 0.0 0.0 0.0
...      ...    ...    ...    ...    ...    ...
32450  9997405307 0.0 1.0 0.0 0.0 0.0 0.0 0.0 0.0
32451  9997522052 0.0 0.0 1.0 0.0 0.0 0.0 0.0 0.0
32452  9999364497 0.0 0.0 1.0 0.0 0.0 0.0 0.0 0.0
32453  B00007MF56 0.0 1.0 0.0 0.0 0.0 0.0 0.0 0.0
32454  B0000DAPP1 0.0 0.0 1.0 0.0 0.0 0.0 0.0 0.0

      Rating(18-29) ... Rating(1920-1930) Rating(1930-1940) \
0      9.0 ... 0.0 0.0
1      0.0 ... 0.0 0.0
2      0.0 ... 0.0 0.0
3      0.0 ... 0.0 0.0

```

4	0.0	...	0.0	0.0
...
32450	0.0	...	0.0	0.0
32451	0.0	...	0.0	0.0
32452	0.0	...	0.0	0.0
32453	9.0	...	0.0	0.0
32454	0.0	...	0.0	0.0

	Rating(1940-1950)	Rating(1950-1960)	Rating(1960-1970)	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	
...	
32450	0.0	0.0	0.0	
32451	0.0	0.0	0.0	
32452	0.0	0.0	0.0	
32453	0.0	0.0	0.0	
32454	0.0	0.0	0.0	

	Rating(1970-1980)	Rating(1980-1990)	Rating(1990-2000)	\
0	0.0	0.0	9.0	
1	0.0	0.0	10.0	
2	0.0	0.0	0.0	
3	0.0	0.0	10.0	
4	0.0	0.0	8.0	
...	
32450	0.0	0.0	0.0	
32451	0.0	0.0	0.0	
32452	0.0	0.0	0.0	
32453	0.0	0.0	0.0	
32454	0.0	0.0	0.0	

	Rating(2000-2010)	Rating(2010-2020)
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0
...
32450	0.0	0.0
32451	0.0	0.0
32452	0.0	0.0
32453	9.0	0.0
32454	5.0	0.0

[32455 rows x 35 columns]

Calculate similarity

```
[125]: product_vec = product_vec.set_index('ISBN')

product_vec = product_vec.dropna(subset = ['1-17'])
```

This is our automatic function to recommend product

```
[126]: def recommend_decade(year_range,ISBN,use_all_year = False):
    neigh = NearestNeighbors(metric='cosine', algorithm='brute', n_neighbors=6)
    book = []
    if use_all_year == False:
        vec = product_vec.loc[product_vec[year_range]== 1,'1-17':
→'Rating(2010-2020)']
        if vec.shape[0] <6:
            print('No enough Books to Recommend!')
        else:
            neigh.fit(vec)
            book = neigh.kneighbors(vec, return_distance = False)
            book = pd.DataFrame(book,index = vec.index)
            book.columns = ['Top1','Top2','Top3','Top4','Top5','Top6']
            book['Top1'] = book.index[book.Top1]
            book['Top2'] = book.index[book.Top2]
            book['Top3'] = book.index[book.Top3]
            book['Top4'] = book.index[book.Top4]
            book['Top5'] = book.index[book.Top5]
            book['Top6'] = book.index[book.Top6]
            result = book[book.index== ISBN]
    if use_all_year == True:
        vec = product_vec
        neigh.fit(vec)
        book = neigh.kneighbors(vec, return_distance = False)
        book = pd.DataFrame(book,index = vec.index)
        book.columns = ['Top1','Top2','Top3','Top4','Top5','Top6']
        book['Top1'] = book.index[book.Top1]
        book['Top2'] = book.index[book.Top2]
        book['Top3'] = book.index[book.Top3]
        book['Top4'] = book.index[book.Top4]
        book['Top5'] = book.index[book.Top5]
        book['Top6'] = book.index[book.Top6]
        result = book[book.index== ISBN]
    return result
```

Recommend books within a specific decade

```
[127]: recommend_decade('2000-2010','0002005018',use_all_year = False)
```

```
[127]:
```

	Top1	Top2	Top3	Top4	Top5	\
ISBN						
0002005018	207042314X	0374253536	2070423530	0671027573	0826452310	

```

Top6
ISBN
0002005018 047121888X

```

```
### Recommend books using all data
```

```
[128]: recommend_decade('2000-2010', '0002005018', use_all_year = True)
```

```
[128]:
```

	Top1	Top2	Top3	Top4	Top5	\
ISBN						
0002005018	0875969313	157765692X	0786014652	0786014512	0786014881	

```

Top6
ISBN
0002005018 0786015233

```

```
[ ]:
```